



Sandia
National
Laboratories

Introduction to MPI- Distributed Computing with GPUs.

Simon Garcia de Gonzalo, Sandia National Laboratories
(slides from Dr. Lena Oden, FernUniversität in Hagen)



Sandia National Laboratories is a multimission laboratory managed and operated by National Technology & Engineering Solutions of Sandia, LLC, a wholly owned subsidiary of Honeywell International Inc., for the U.S. Department of Energy's National Nuclear Security Administration under contract DE-NA0003525.

Multi-GPU Computing: What you will learn



- CUDA-aware MPI
- Example: Jacobi-Solver
- Under the hood (why you should use CUDA-aware MPI)
 - GPUs in Clusters
 - CUDA Unified Virtual Addressing
 - GPUDirect P2P and GPUDirect RDMA

Message Passing Interface - MPI



- Standard to exchange data between processes via messages
 - Defines API to exchange messages
 - Point to Point: e.g. MPI_Send, MPI_Recv
 - Collectives: e.g. MPI_Reduce, MPI_Allreduce, MPI_Bcast
- Multiple implementations (open source and commercial)
 - Bindings for C/C++, Fortran, Python, ...
 - e.g. MPICH, OpenMPI, MVAPICH, IBM Spectrum MPI, Cray MPT, ParaStation MPI, ...

Example MPI Program

```
#include <mpi.h>
#include <stdio.h>

int main(int argc, char** argv) {
    // Initialize the MPI environment
    MPI_Init(NULL, NULL);
    int size;
    MPI_Comm_size(MPI_COMM_WORLD, &size);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    MPI_Finalize();
}
```

```
mpicc -o hello_mpi.out hello_mpi.c
mpirun -n 4 ./hello_mpi.out
```




CUDA-aware MPI

CUDA-aware MPI allows you to use Pointers to GPU-Memory as source and destination

```
//MPI rank 0
MPI_Send(s_buf_d, n, MPI_BYTE, size-1, tag, MPI_COMM_WORLD);

//MPI size-1
MPI_Recv(r_buf_d, n, MPI_BYTE, 0, tag, MPI_COMM_WORLD, MPI_STATUS_IGNORE);
```



Pointer to
GPU
memory!



- Launch one process per GPU
- How to use CUDA-aware MPI
 - *MVAPICH*: `$MV2_USE_CUDA=1 mpirun -np ${np} ./myapp <args>`
 - *Open MPI*: *CUDA-aware features are enabled per default (using UCX)*
 - *Cray*: `MPICH_RDMA_ENABLED_CUDA`
 - *IBM Spectrum MPI*: `$mpirun -gpu -np ${np} ./myapp <args>`
 - *ParaStation MPI* (using UCX): `$PSP_CUDA=1 mpirun -np ${np} ./myapp <args>`
- On JUWELS Booster:
 - Load CUDA-aware OpenMPI or ParaStation MPI modules
 - GPU-tuning is done via module loading
 - `srun --gres=gpu:4 -n ${np} ./myapp <args>`

How to compile




```
nvcc -o my_kernel.o $(NVCC_FLAGS) my_kernels.cu -c  
mpicc -o my_multiGPUapp -lcudart my_kernel.o my_multiGPUapp.c
```

Example: Jacobi Solver



- Solves the 2D-Laplace Equation on a rectangle
- $\Delta u(x, y) = 0 \forall (x, y) \in \Omega \setminus \delta\Omega$
- Dirichlet boundary conditions on left/right boundaries (constant values on boundaries)
- Reflecting boundaries on top and bottom
- Iterative solver:
 $u(t+1) = f(u(t))$




Example: Jacobi solver



While not converged

do Jacobi step:

```
int iy = blockIdx.y * blockDim.y + threadIdx.y
int ix = blockIdx.x * blockDim.x + threadIdx.x
if (iy < ny-1 && ix < nx-1) {
    new_val =
        0.25 * (a[iy * nx + ix + 1] + a[iy * nx + ix - 1
                                             a[(iy + 1) * nx + ix] + a[(iy - 1) * nx + ix]);
    a_new[iy * nx + ix] = new_val;
```




apply boundary Condition



swap a_new and a

next iteration

Domain Decomposition



Minimize number of neighbors:
Communicate to less neighbors
Optimal for latency bound communication,
Continuous Transfers



Minimize surface area/volume ratio:
Communicate less data
Optimal for bandwidth bound communication
Non-Continuous Transfers


Jacobi example: Top and Bottom Boundaries

```
MPI_Sendrecv(a_new_d+offset_last_row, m-2, MPI_DOUBLE,  
            b_nb, 1, a_new_d+offset_top_boundary, m-2,  
            MPI_DOUBLE, t_nb, 1, MPI_COMM_WORLD,  
            MPI_STATUS_IGNORE);
```

bottom
neighbor

```
MPI_Sendrecv(a_new_d+offset_first_row, m-2, MPI_DOUBLE,  
            t_nb, 0, a_new_d+offset_bottom_boundary, m-2,  
            MPI_DOUBLE, b_nb, 0, MPI_COMM_WORLD,  
            MPI_STATUS_IGNORE);
```



top
neighbor



Jacobi: Left (and right) neighbours




```
//right neighbor omitted
pack<<<gs,bs,0,s>>>(to_left_d, u_new_d, n, m);
cudaStreamSynchronize(s);
MPI_Sendrecv(to_left_d, n-2, MPI_DOUBLE, l_nb, 0,
             from_right_d, n-2, MPI_DOUBLE, r_nb, 0,
             MPI_COMM_WORLD, MPI_STATUS_IGNORE );
unpack<<<gs,bs,0,s>>>(u_new_d, from_right_d, n, m);
cudaStreamSynchronize(s);
```



Process Mapping on Multi GPU Systems

One GPU per Process



Distribute GPUs to local Nodes




```
MPI_Comm local_comm;  
MPI_Comm_split_type(MPI_COMM_WORLD, MPI_COMM_TYPE_SHARED, rank,  
                     MPI_INFO_NULL, &local_comm);  
int local_rank = -1;  
MPI_Comm_rank(local_comm, &local_rank);  
MPI_Comm_free(&local_comm);  
  
int num_devs = 0;  
cudaGetDeviceCount(&num_devs);  
cudaSetDevice(local_rank%num_devs);
```

Needed, if
resource
manager handles
GPU-affinity

Process Mapping on Multi GPU Systems


One GPU per Process




CUDA Unified Virtual Addressing




- One address space for all CPU and GPU memory
 - Determine physical memory location from a pointer value
 - Enable libraries to simplify their interfaces (e.g. MPI)
- Supported on devices with compute capability 2.0+ for
 - 64-bit applications on Linux and Windows (+TCC)



Basics: GPUDirect




CUDA-aware MPI with GPUDirect RDMA




```
MPI_Send(s_buf_d, size, MPI_BYTE, 1, tag, MPI_COMM_WORLD);  
MPI_Recv(r_buf_d, size, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &stat);
```

CUDA-aware MPI without GPUDirect RDMA



```
MPI_Send(s_buf_d, size, MPI_BYTE, 1, tag, MPI_COMM_WORLD);  
MPI_Recv(r_buf_d, size, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &stat);
```

Regular MPI




```
cudaMemcpy(s_buf_h, s_buf_d, size, cudaMemcpyDeviceToHost);  
MPI_Send(s_buf_h, size, MPI_BYTE, 1, tag, MPI_COMM_WORLD);  
  
MPI_Recv(r_buf_h, size, MPI_BYTE, 0, tag, MPI_COMM_WORLD, &stat);  
cudaMemcpy(r_buf_d, r_buf_h, size, cudaMemcpyHostToDevice);
```

Performance Results GPUDirect RDMA




OpenMPI 4.1.0RC1 + UCX 1.9.0 on JUWELS Booster



Performance Results GPUDirect P2P



OpenMPI 4.1.0RC1 + UCX 1.9.0 on JUWELS-Booster



Summary



- CUDA-aware MPI allows efficient communication for multi-GPU applications
 - Allows MPI-communication operations from GPU memory buffers
 - Simplified programming
 - Use GPUDirect technologies for performance
 - Minimizes data copies
- Most MPI versions have support for CUDA-aware MPI
- Best practice: One process per rank
 - Use local communicator in MPI