




Fakultät für
**Mathematik und
Informatik**

CUDA Graphs and Device-initiated Communication with NVSHMEM

Lena Oden, FernUniversität in Hagen

Jiri Kraus, Principal Devtech Compute, Nvidia

Multi GPU Jacobi Nsight Systems Timeline



Asynchronous Task Graph


A Graph Node Is A CUDA Operation

Sequence of operations (nodes), connected by dependencies

Operations are one of:


| | |
|---------------------|----------------------------|
| Kernel Launch | CUDA kernel running on GPU |
| CPU Function Call | Callback function on CPU |
| Memcpy/Memset | GPU data management |
| Mem Alloc/Free | Memory management |
| External Dependency | External semaphores/events |
| Sub-Graph | Graphs are hierarchical |

Nodes within a graph can also span multiple devices



Three-Stage Execution Model


Define



Single Graph “Template”

Created in host code
or built up from libraries

Instantiate




Multiple “Executable Graphs”

Snapshot of templates

Sets up & initializes GPU execution structures
(create once, run many times)

Execute



Executable Graphs Running in CUDA Streams


Concurrency in graph **is not** limited by stream

Where is performance coming from?

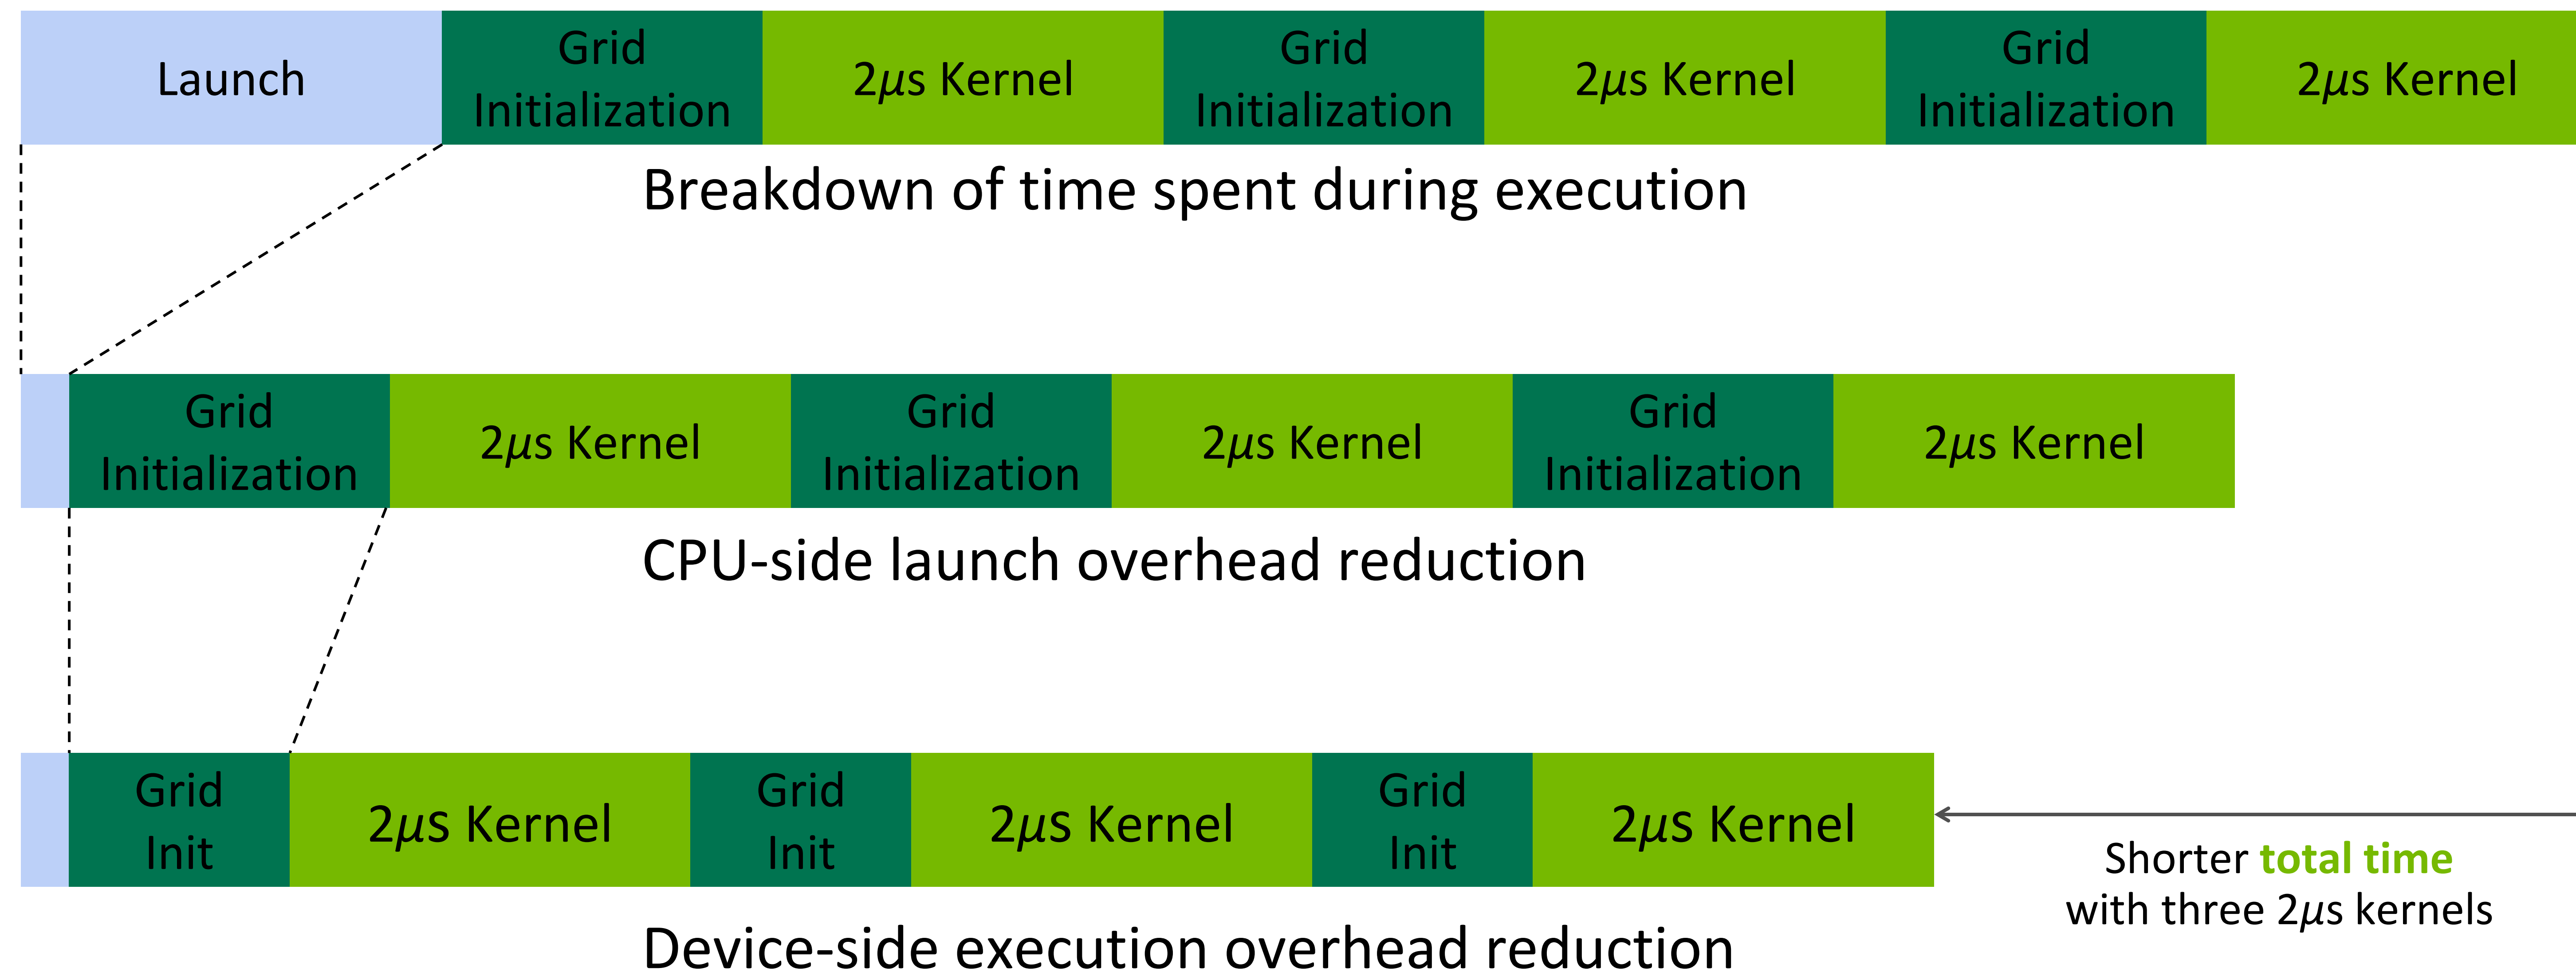


Breakdown of time spent during execution

Where is performance coming from?



Where is performance coming from?



Capture Stream work into a Graph Create A Graph With Two Lines of Code

```
{  
  
    cudaStreamBeginCapture(compute_stream, cudaStreamCaptureModeGlobal);  
  
    cudaMemsetAsync(l2_norm_d, 0, sizeof(real), compute_stream);  
    cudaEventRecord(reset_l2norm_done, compute_stream);  
  
    ...  
  
    cudaStreamWaitEvent(compute_stream, push_done, 0);  
  
    cudaStreamEndCapture(compute_stream, graphs[calculate_norm]+is_even);  
  
    std::swap(a_new, a);  
    iter++;  
}
```


Capture Stream work into a Graph

Create A Graph With Two Lines of Code

```
{  
  
    cudaStreamBeginCapture(compute_stream, cudaStreamCaptureModeGlobal);  
  
    cudaMemsetAsync(l2_norm_d, 0, sizeof(real), compute_stream);  
    cudaEventRecord(reset_l2norm_done, compute_stream);  
  
    ...  
  
    cudaStreamWaitEvent(compute_stream, push_done, 0);  
  
    cudaStreamEndCapture(compute_stream, graphs[calculate_norm]+is_even);  
  
    std::swap(a_new, a);  
    iter++;  
}
```

CUDA Graph Management API

Instantiate CUDA Graphs

```
__host__ cudaError_t cudaGraphInstantiateWithFlags ( cudaGraphExec_t* pGraphExec, cudaGraph_t graph,  
                                                    unsigned long long flags )  
  
__host__ cudaError_t cudaGraphInstantiate ( cudaGraphExec_t* pGraphExec, cudaGraph_t graph,  
                                             cudaGraphNode_t* pErrorNode, char* pLogBuffer, size_t bufferSize )
```


- pGraphExec [OUT] Returns instantiated graph
 - graph [IN]: Graph to instantiate
 - flags [IN]: Flags to control instantiation (cudaGraphInstantiateFlagAutoFreeOnLaunch | **cudaGraphInstantiateFlagUseNodePriority**).
 - pErrorNode [OUT]: In case of an instantiation error, this may be modified to indicate a node contributing to the error
 - pLogBuffer [OUT]: A character buffer to store diagnostic messages
 - bufferSize [IN]: Size of the log buffer in bytes
- Returns: cudaSuccess, cudaErrorInvalidValue

New Execution Mechanism Graphs Can Be Generated Once Then Launched Repeatedly

```
while (l2_norm > tol && iter < iter_max) {

    cudaGraphLaunch(graph_calc_norm_exec[iter%2],
                    compute_stream);
    cudaStreamSynchronize(compute_stream);
    MPI_Allreduce(l2_norm_h, &l2_norm, 1,
                  MPI_REAL_TYPE, MPI_SUM,
                  MPI_COMM_WORLD);
    l2_norm = std::sqrt(l2_norm);

    if (!csv && 0 == rank && (iter % 100) == 0) {
        printf("%5d, %0.6f\n", iter, l2_norm);
    }
}
```



```
cudaGraphDebugDotPrint(graphs[calculate_norm][0],
                        "jacobi_graph.dot", 0)
```

```
dot -Tpng jacobi_graph.dot -o jacobi_grap.png
```


CUDA Graph Management API

```
__host__ cudaError_t cudaGraphDestroy( cudaGraph_t graph)
```

- graph [IN]: Graph to destroy

Returns: cudaSuccess, cudaErrorInvalidValue

Destroys the graph specified by graph, as well as all of its nodes.


```
__host__ cudaError_t cudaGraphExecDestroy ( cudaGraphExec_t graphExec )
```

- graphExec [IN]: Executable graph to destroy

Returns: cudaSuccess, cudaErrorInvalidValue

Destroys the executable graph specified by graphExec.

Multi GPU Jacobi Nsight Systems Timeline



CPU-Initiated Communication

Compute on GPU

Communication from CPU

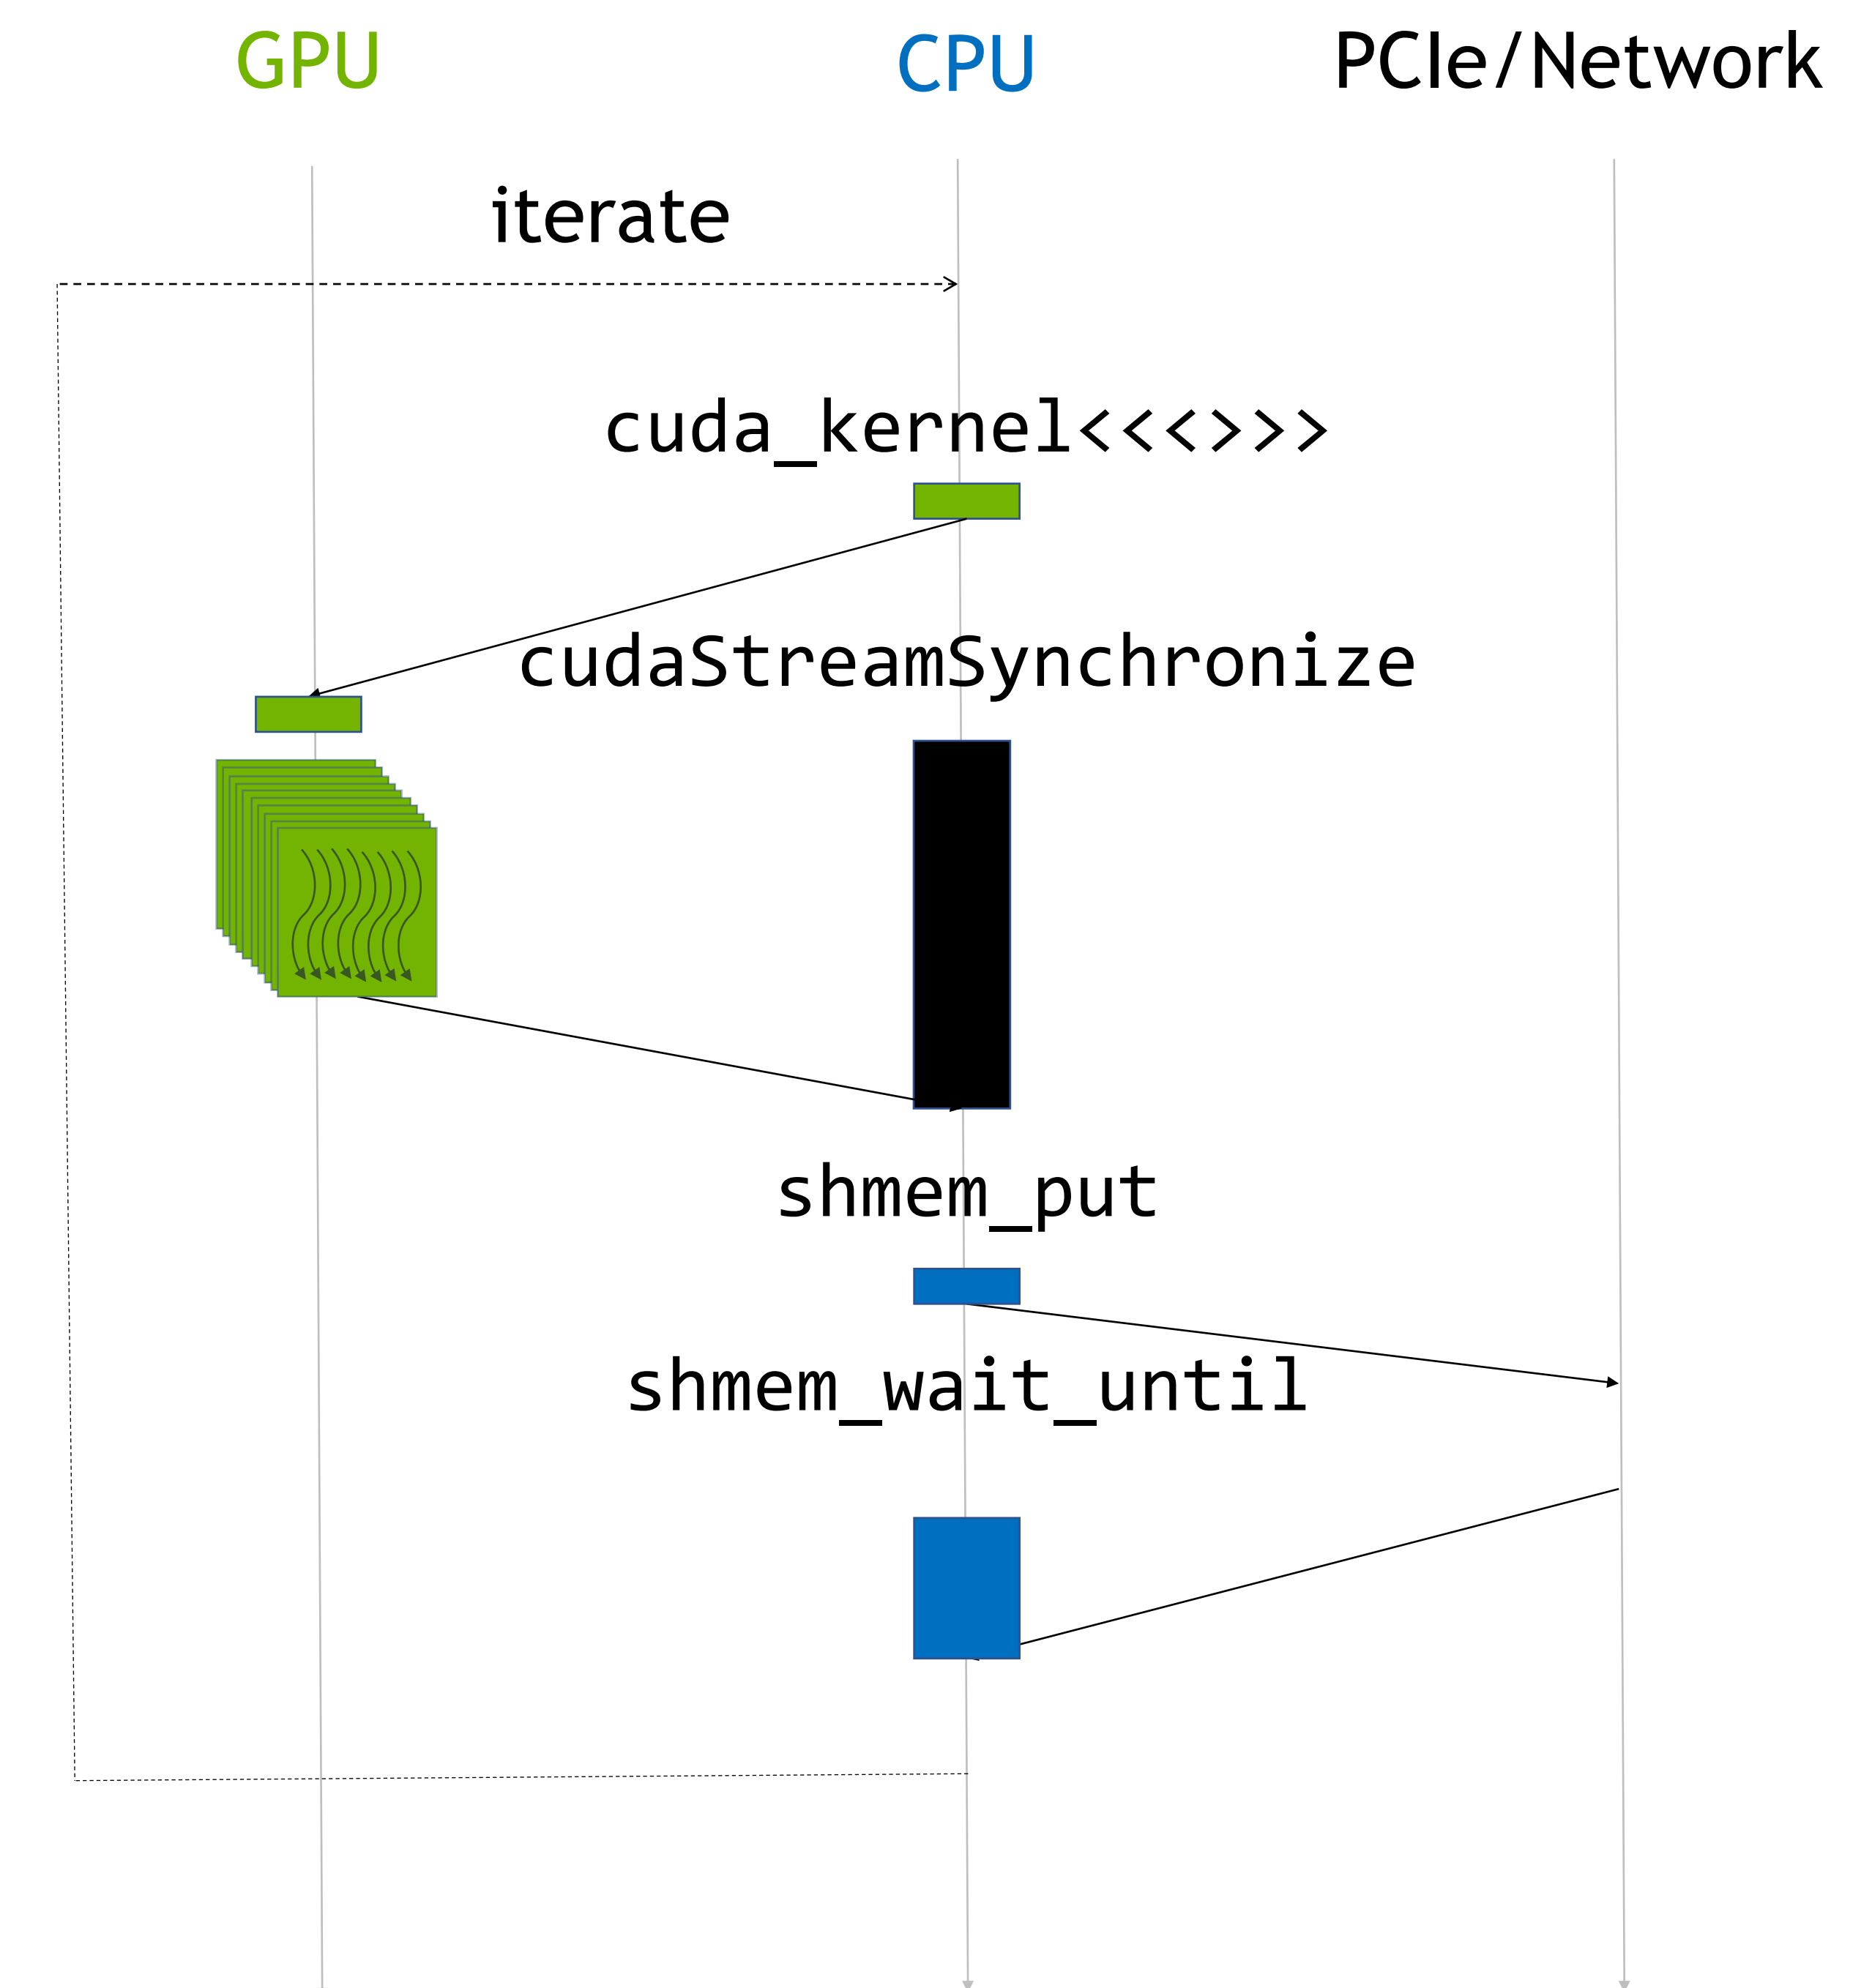
Synchronization at boundaries

Commonly used model, but –

Offload latencies in critical path

Communication is not overlapped

Hiding increased code complexity, not hiding limits strong scaling



GPU-Initiated Communication

Compute on GPU

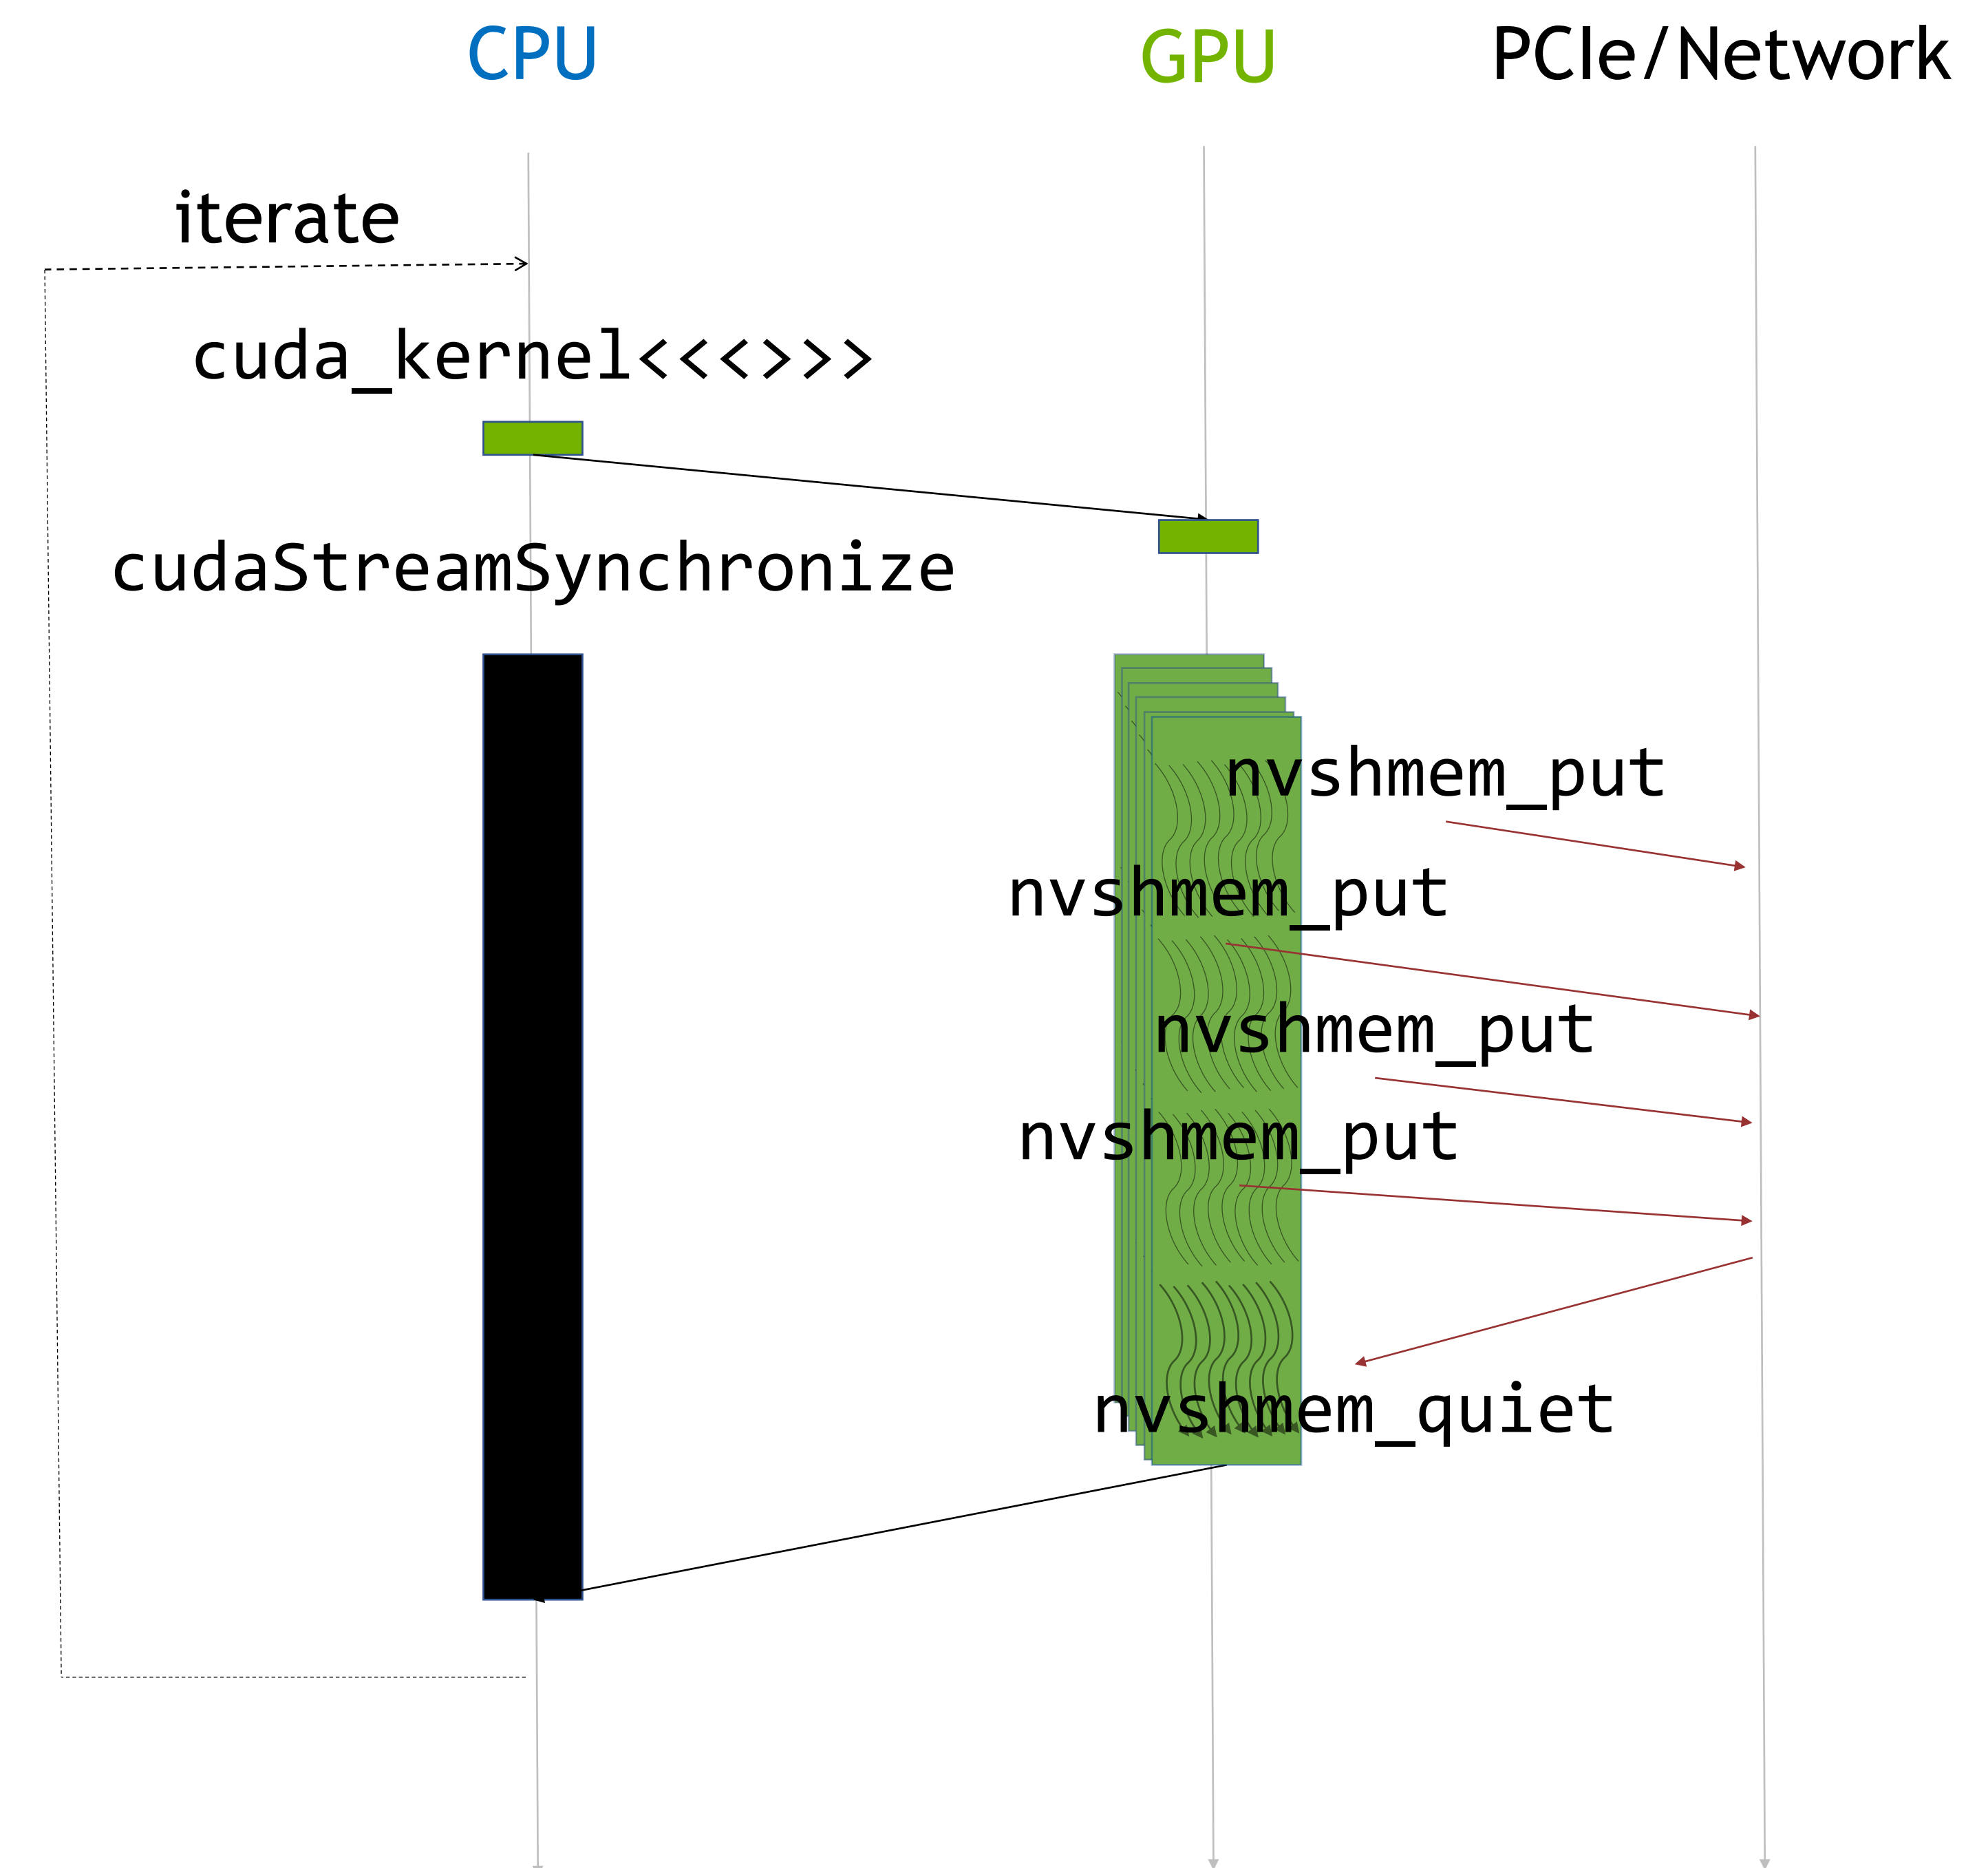
Communication from GPU

Benefits

Eliminates offloads latencies

Compute and communication overlap by threading

Easier to express algorithms with inline communication




Thread-level communication

- Allows fine grained communication and computation overlap
- Efficient mapping to NVLink fabric on DGX systems

```
__global__ void stencil_single_step(float *u, float *v, ...) {
    int ix = get_ix(blockIdx, blockDim, threadIdx);
    int iy = get_iy(blockIdx, blockDim, threadIdx);
    compute(u, v, ix, iy);
    // Thread-level data communication API

    if (iy == 1)
        nvshmem_float_p(u+(ny+1)*nx+ix, u[nx+ix], top_pe);
    if (iy == ny)
        nvshmem_float_p(u+ix, u[ny*nx+ix], bottom_pe);
}

for (int iter = 0; iter < N; iter++) {
    swap(u, v);
    stencil_single_step<<<..., stream>>>(u, v, ...);
    nvshmem_barrier_all_on_stream(stream);
}
```

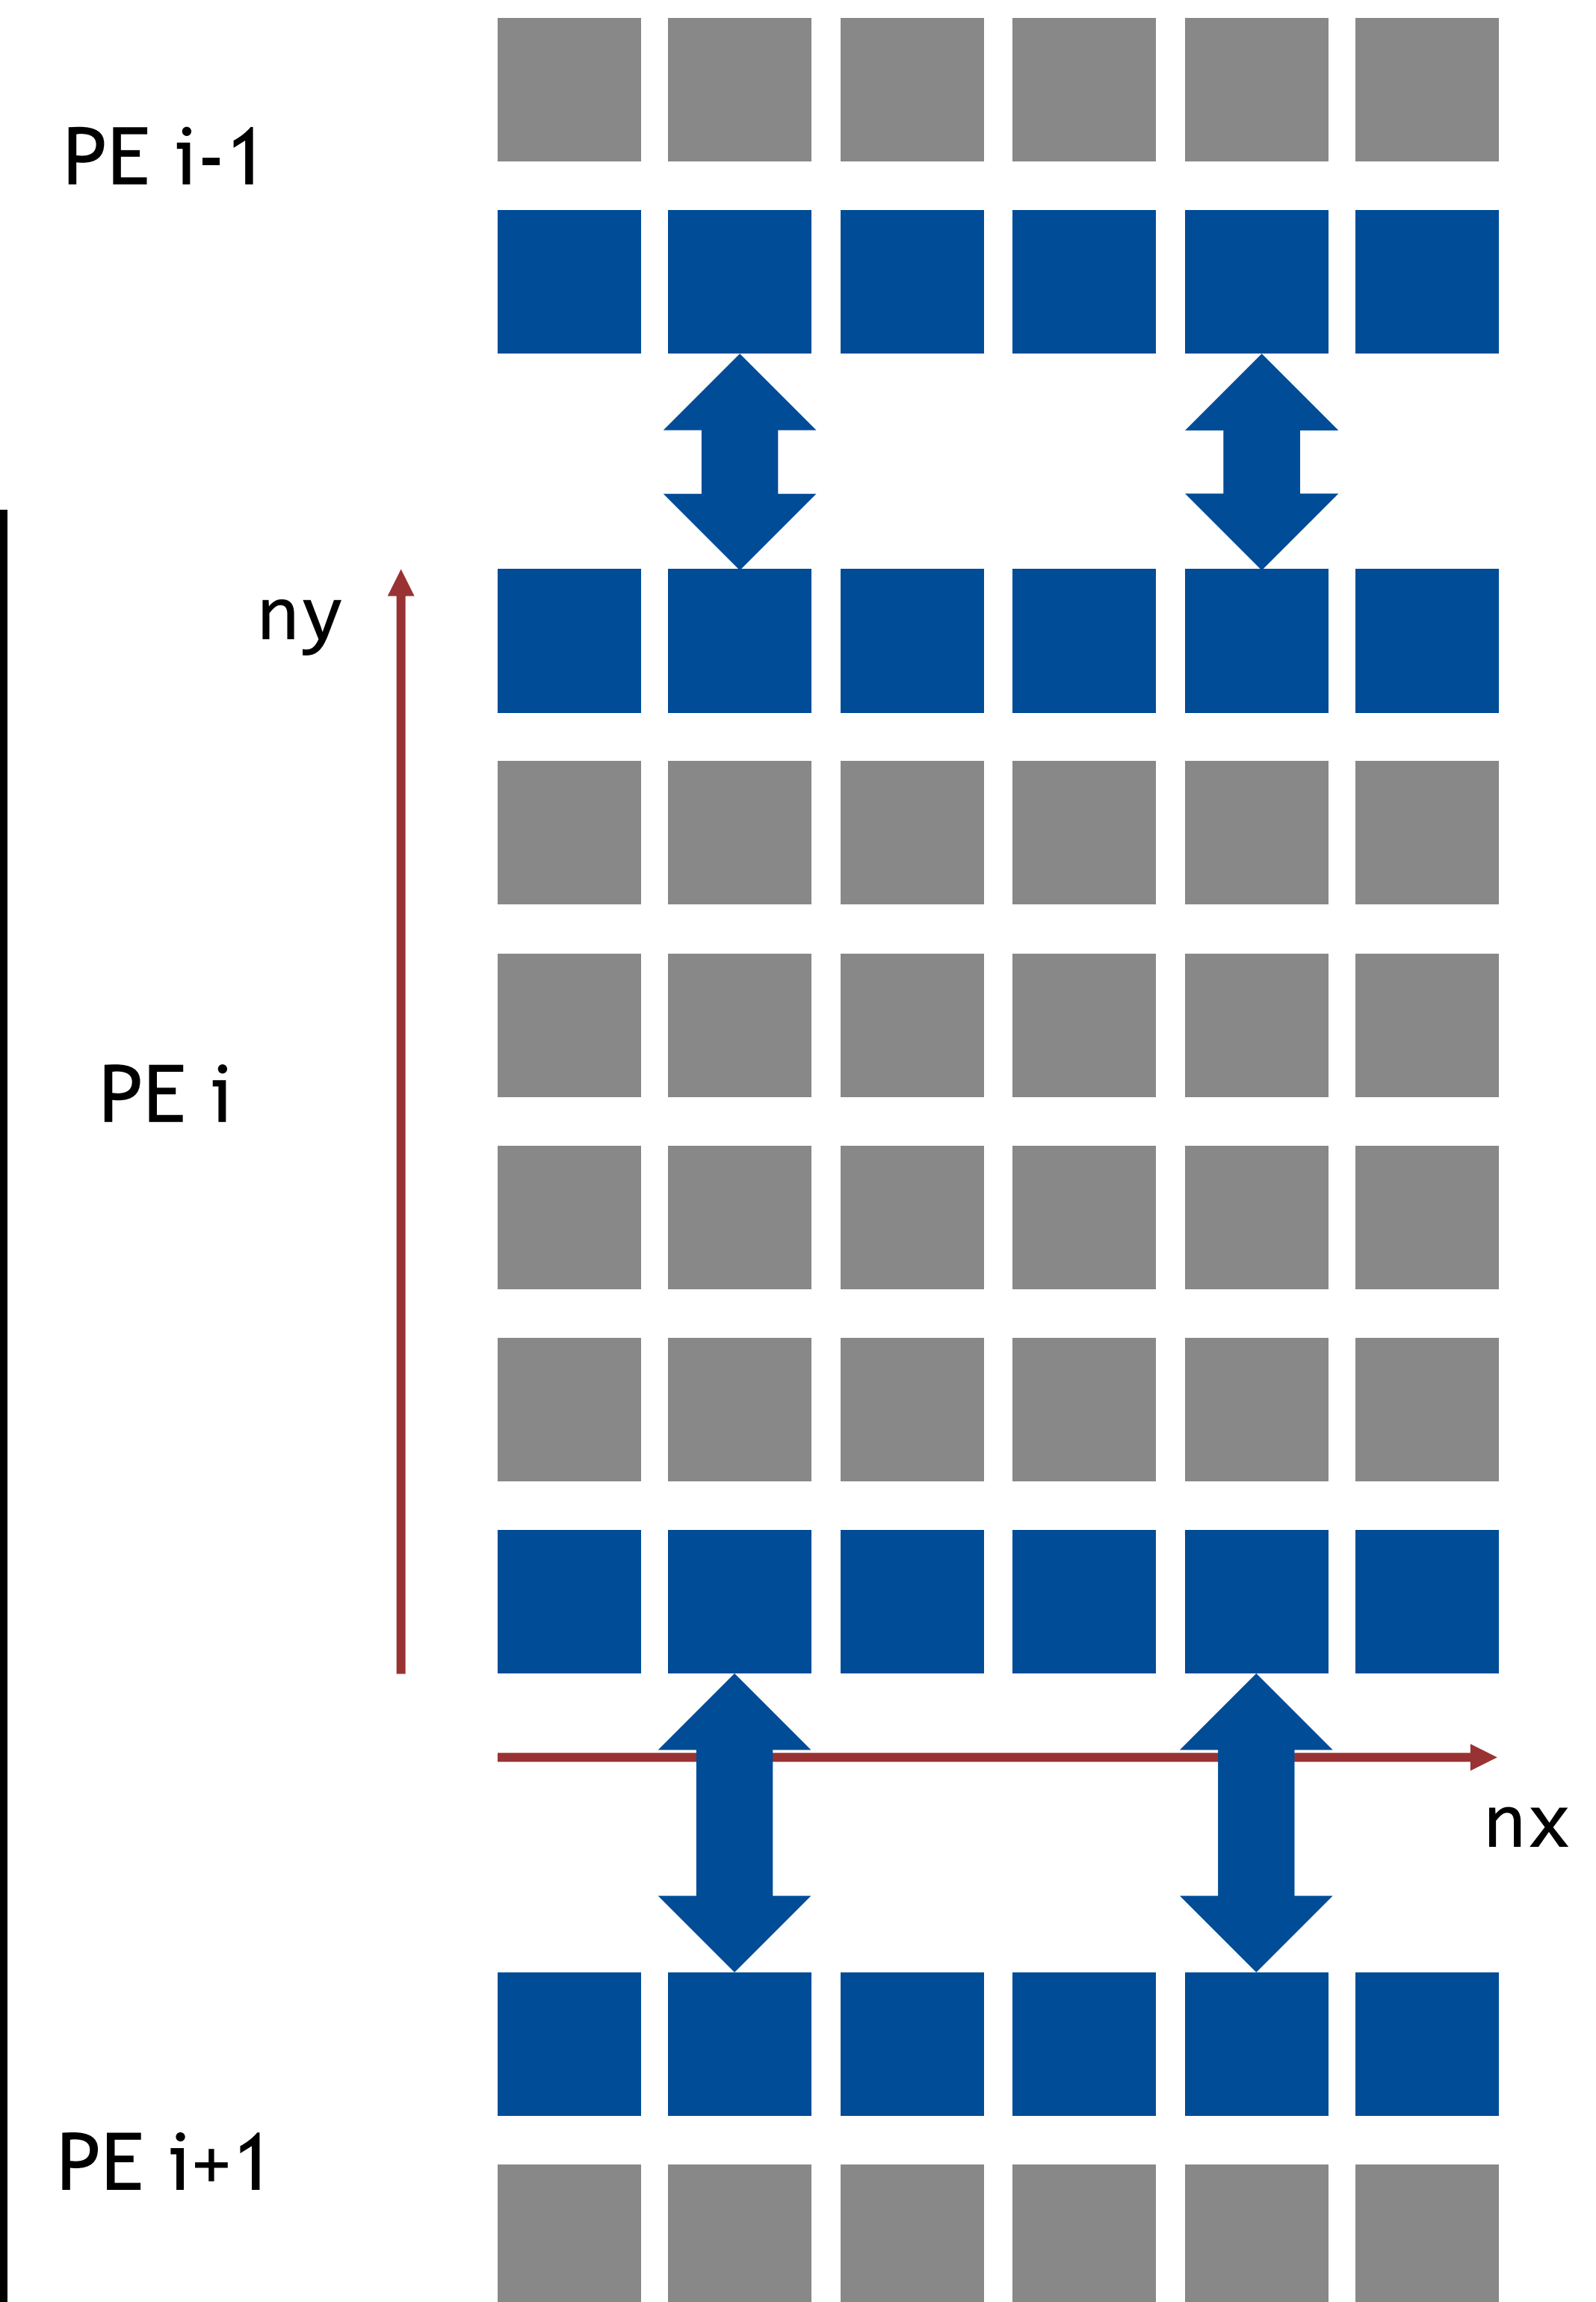


Thread-GROUP communication

- More efficient data transfers over networks like IB
- NVSHMEM operations can be issued by all threads in a block/warp
- Still allows inter-warp/inter-block overlap

```
__global__ void stencil_single_step(float * u, float *v, ...) {
    int ix = get_ix(blockIdx.x, blockDim.x, threadIdx.x);
    int iy = get_iy(blockIdx.y, blockDim.y, threadIdx.y);
    compute(u, v, ix, iy);
    // Thread block-level communication API
    int boffset = get_block_offset(blockIdx.x, blockDim.x);
    if (blockIdx.y == 0)
        nvshmemx_float_put_nbi_block(u+(ny+1)*nx+boffset, u+nx+boffset, blockDim.x, top_pe);
    if (blockIdx.y == (blockDim.y-1))
        nvshmemx_float_put_nbi_block(u+boffset, u+ny*nx+boffset, blockDim.x, bottom_pe);
}

for (int iter = 0; iter < N; iter++) {
    swap(u, v);
    stencil_single_step<<<..., stream>>>(u, v, ...);
    nvshmem_barrier_all_on_stream(stream);
}
```

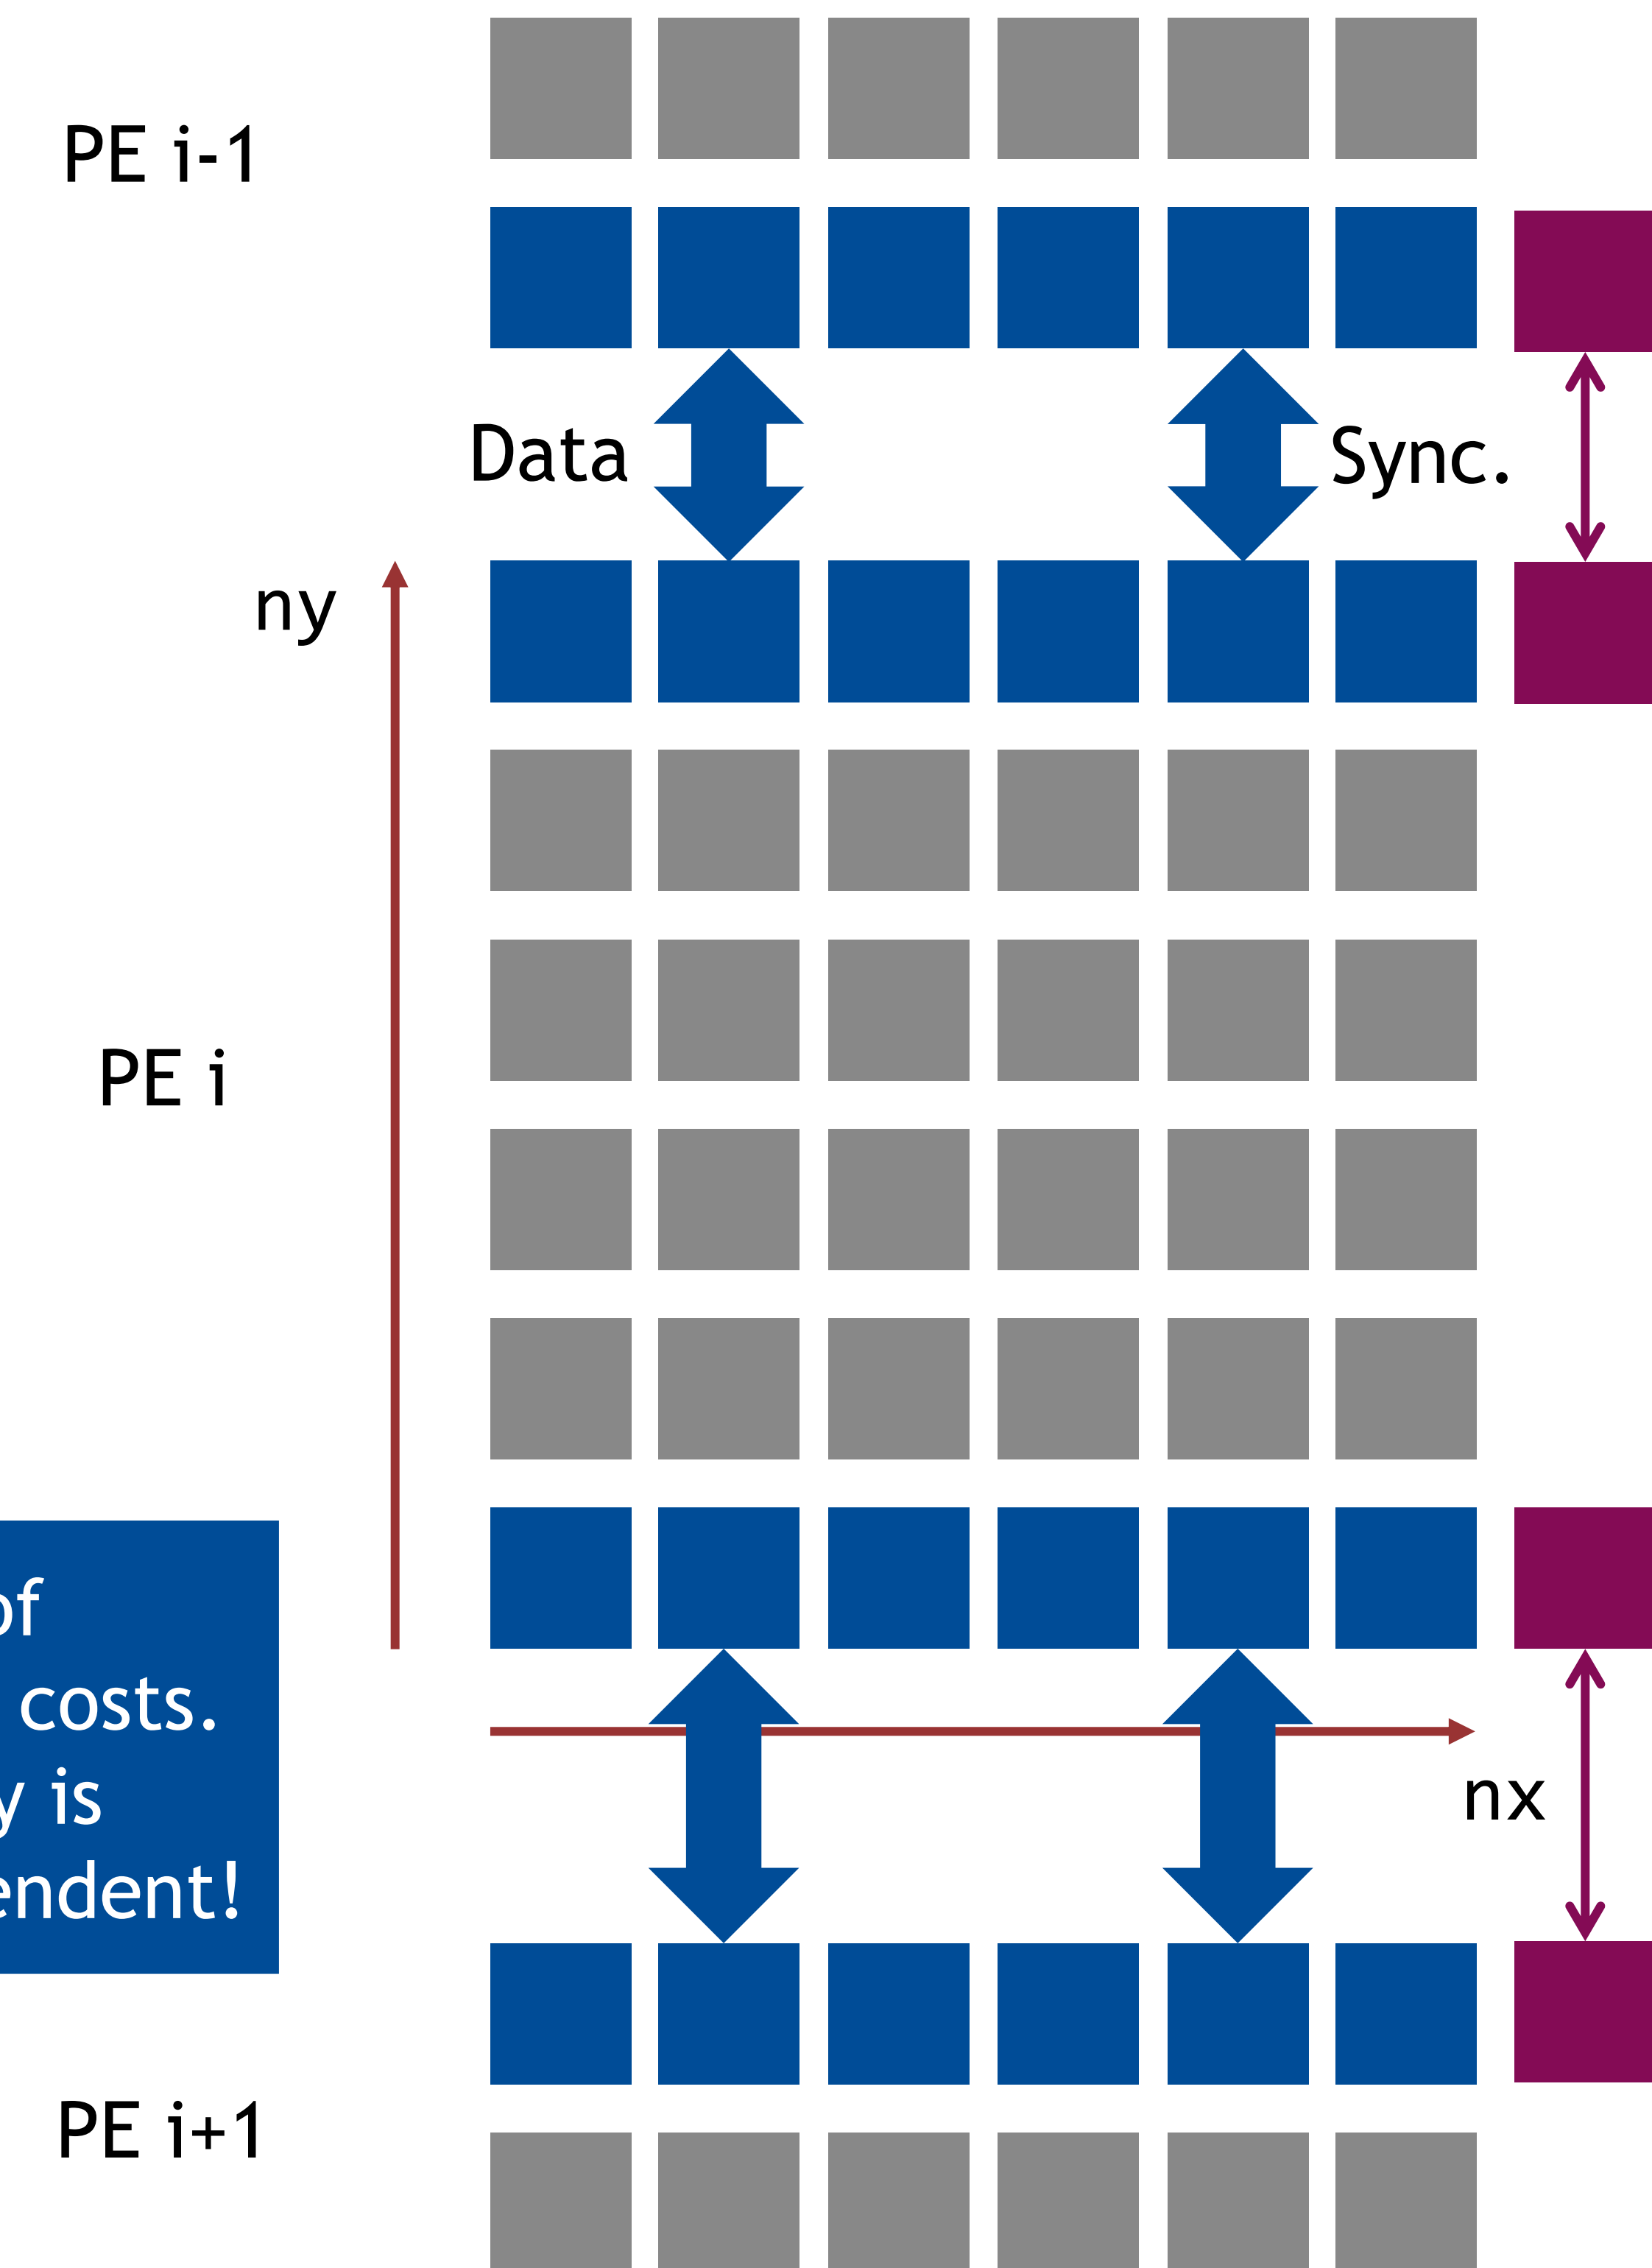


In-Kernel Synchronization

- Point-to-point synchronization across PEs within a kernel
- Enables kernel fusion

```
__global__ void stencil_multi_step(float *u, float *v, int N, int *sync, ...) {
    int ix = get_ix(blockIdx blockIdx.threadIdx);
    int iy = get_iy(blockIdx blockIdx.threadIdx);
    for (int iter = 0; iter < N; iter++) {
        swap(u, v); compute(u, v, ix, iy);
        // Thread block-level data exchange (assume even/odd iter buffering)
        int boffset = get_block_offset(blockIdx blockIdx);
        if (blockIdx.y == 0)
            nvshmemx_float_put_nbi_block(u+(ny+1)*nx+boffset, u+nx+boffset, blockDim.x, top_pe);
        if (blockIdx.y == (blockDim.y-1))
            nvshmemx_float_put_nbi_block(u + boffset, u+ny*nx+boffset, blockDim.x, bottom_pe);
        if (blockIdx.y == 0 || blockIdx.y == (blockDim.y-1)) {
            __syncthreads();
            nvshmem_quiet();
            if (threadIdx.x == 0 && threadIdx.y == 0) {
                nvshmem_atomic_inc(sync, top_pe);
                nvshmem_atomic_inc(sync, bottom_pe);
            }
        }
        nvshmem_wait_until(sync, NVSHMEM_CMP_GT, 2*iter*gridDim.x);
    }
}
```

Be aware of
synchronization costs.
Best strategy is
application dependent!



Collective Kernel Launch

Ensures progress when using device-side inter-kernel synchronization

| NVSHMEM Usage | CUDA Kernel launch |
|----------------------------------|--|
| Device-Initiated Communication | Execution config syntax <<<...>>> or launch APIs |
| Device-Initiated Synchronization | <code>nvshmemx_collective_launch</code> |

- CUDA's throughput computing model allows (encourages) grids much larger than a GPU can fit
- Inter-kernel synchronization requires producer and consumer threads to execute concurrently
- Collective launch guarantees co-residency using CUDA cooperative launch and requirement of 1PE/GPU


```
__device__ void nvshmem_TYPENAME_p(TYPE *dest, TYPE value, int pe)
```

- dest [OUT]: Symmetric address of the destination data object.
- value [IN]: The value to be transferred to dest.
- pe [IN]: The number of the remote PE.

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html#nvshmem-p>

TYPENAME can be: float, double, char, schar, short, int, long, longlong, uchar, ushort, uint, ulong, ulonglong, ..., ptrdiff

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html#stdrmatypes>




Thread-level communication

- Allows fine grained communication and computation overlap
- Efficient mapping to NVLink fabric on DGX systems

```
__global__ void stencil_single_step(float *u, float *v, ...) {
    int ix = get_ix(blockIdx, blockDim, threadIdx);
    int iy = get_iy(blockIdx, blockDim, threadIdx);
    compute(u, v, ix, iy);
    // Thread-level data communication API

    if (iy == 1)
        nvshmem_float_p(u+(ny+1)*nx+ix, u[nx+ix], top_pe);
    if (iy == ny)
        nvshmem_float_p(u+ix, u[ny*nx+ix], bottom_pe);
}

for (int iter = 0; iter < N; iter++) {
    swap(u, v);
    stencil_single_step<<<..., stream>>>(u, v, ...);
    nvshmem_barrier_all_on_stream(stream);
}
```




```
__device__ void nvshmemx_TYPENAME_put_nbi_block(TYPE *dest, const TYPE *source, size_t nelems, int pe)
```

- dest [OUT]: Symmetric address of the destination data object.
- source [IN]: Symmetric address of the object containing the data to be copied.
- nelems [IN]: Number of elements in the dest and source arrays.
- pe [IN]: The number of the remote PE.

Cooperative call: Needs to be called by all threads in a block. thread and warp are also available.

x in nvshmemx marks API as extension of the OpenSHMEM APIs.



https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html?highlight=nvshmemx_typename_put_nbi_block#nvshmem-put-nbi

TYPENAME can be: float, double, char, schar, short, int, long, longlong, uchar, ushort, uint, ulong, ulonglong, ..., ptrdiff

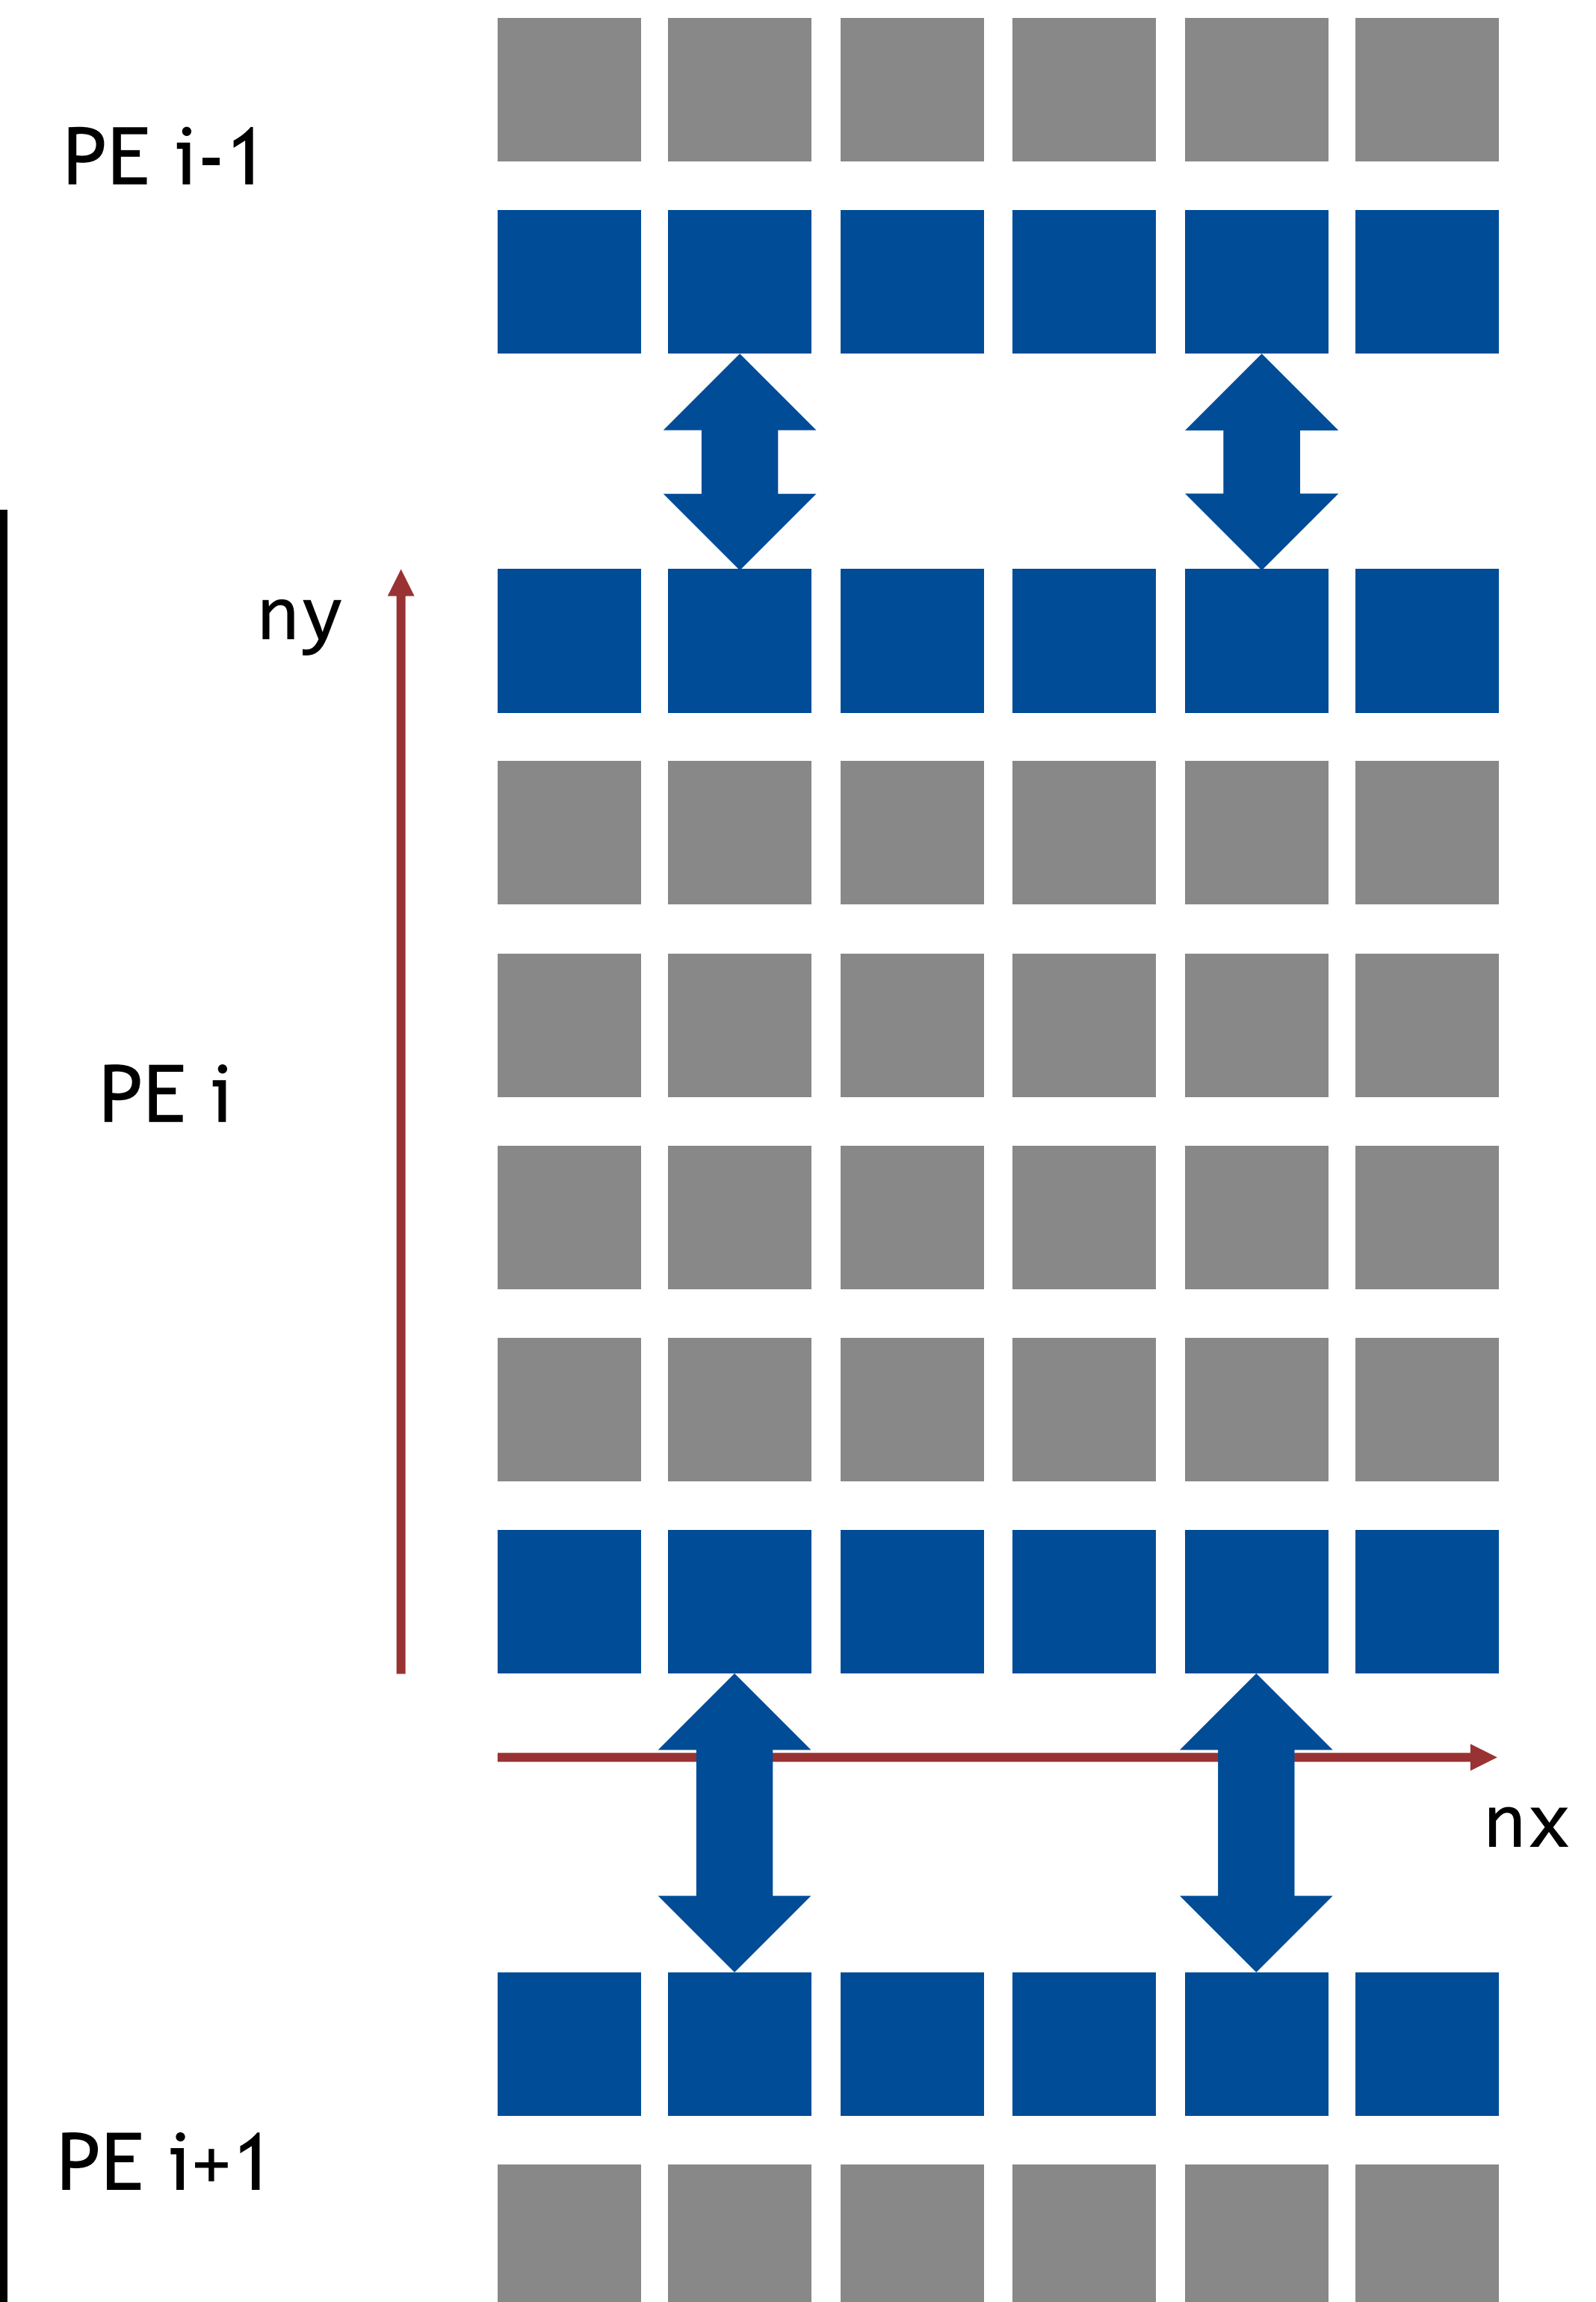
<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html#stdrmatypes>

Thread-GROUP communication

- More efficient data transfers over networks like IB
- NVSHMEM operations can be issued by all threads in a block/warp
- Still allows inter-warp/inter-block overlap

```
__global__ void stencil_single_step(float * u, float *v, ...) {
    int ix = get_ix(blockIdx.x, blockDim.x, threadIdx.x);
    int iy = get_iy(blockIdx.y, blockDim.y, threadIdx.y);
    compute(u, v, ix, iy);
    // Thread block-level communication API
    int boffset = get_block_offset(blockIdx.x, blockDim.x);
    if (blockIdx.y == 0)
        nvshmemx_float_put_nbi_block(u+(ny+1)*nx+boffset, u+nx+boffset, blockDim.x, top_pe);
    if (blockIdx.y == (blockDim.y-1))
        nvshmemx_float_put_nbi_block(u+boffset, u+ny*nx+boffset, blockDim.x, bottom_pe);
}

for (int iter = 0; iter < N; iter++) {
    swap(u, v);
    stencil_single_step<<<..., stream>>>(u, v, ...);
    nvshmem_barrier_all_on_stream(stream);
}
```



```
__device__ void nvshmem_quiet(void)
```

Ensures completion of all operations on symmetric data objects issued by the calling PE.

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/ordering.html#nvshmem-quiet>


```
__device__ inline void nvshmemx_signal_op(uint64_t *sig_addr, uint64_t signal, int sig_op, int pe)
```

- sig_addr [OUT]: Symmetric address of the signal word to be updated.
- signal [IN]: The value used to update sig_addr.
- sig_op [IN]: Operation used to update sig_addr with signal. (NVSHMEM_SIGNAL_SET or NVSHMEM_SIGNAL_ADD)
- pe [IN]: The number of the remote PE.

x in nvshmemx marks API as extension of the OpenSHMEM APIs.

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/signal.html#nvshmemx-signal-op>

NVSHMEM API

```
__device__ void nvshmem_TYPENAME_atomic_inc(TYPE *dest, int pe)
```

- dest [OUT]: Symmetric address of the signal word to be updated.
- pe [IN]: The number of the remote PE.

These routines perform an atomic increment operation on the dest data object on PE.

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/amo.html#nvshmem-atomic-inc>

TYPENAME can be: float, double, char, schar, short, int, long, longlong, uchar, ushort, uint, ulong, ulonglong, ..., ptrdiff

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html#stdrmatypes>

```
__device__ void nvshmem_TYPENAME_wait_until_all(TYPE *ivars, size_t nelems, const int *status,  
                                                int cmp, TYPE cmp_value)  
  
__device__ void nvshmem_TYPENAME_wait_until(TYPE *ivar int cmp, TYPE cmp_value)
```

- `ivars[N]`: Symmetric address of an array of remotely accessible data objects. | Symmetric address of a remotely accessible data object.
- `nelems [IN]`: The number of elements in the `ivars` array.
- `status [IN]`: Local address of an optional mask array of length `nelems` that indicates which elements in `ivars` are excluded from the wait set. Set to `NULL` when not used.
- `cmp [IN]`: A comparison operator (`NVSHMEM_CMP_EQ`, `NVSHMEM_CMP_NE`, `NVSHMEM_CMP_GT`, `NVSHMEM_CMP_GE`, `NVSHMEM_CMP_LT`, `NVSHMEM_CMP_LE`) that compares **elements of `ivars`** `ivar` with `cmp_value`.
- `cmp_value [IN]`: The value to be compared with the objects pointed to by `ivars`.

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/sync.html#nvshmem-wait-until-all>

<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/sync.html#nvshmem-wait-until>

`TYPENAME` can be: `float`, `double`, `char`, `schar`, `short`, `int`, `long`, `longlong`, `uchar`, `ushort`, `uint`, `ulong`, `ulonglong`, ..., `ptrdiff` (see:

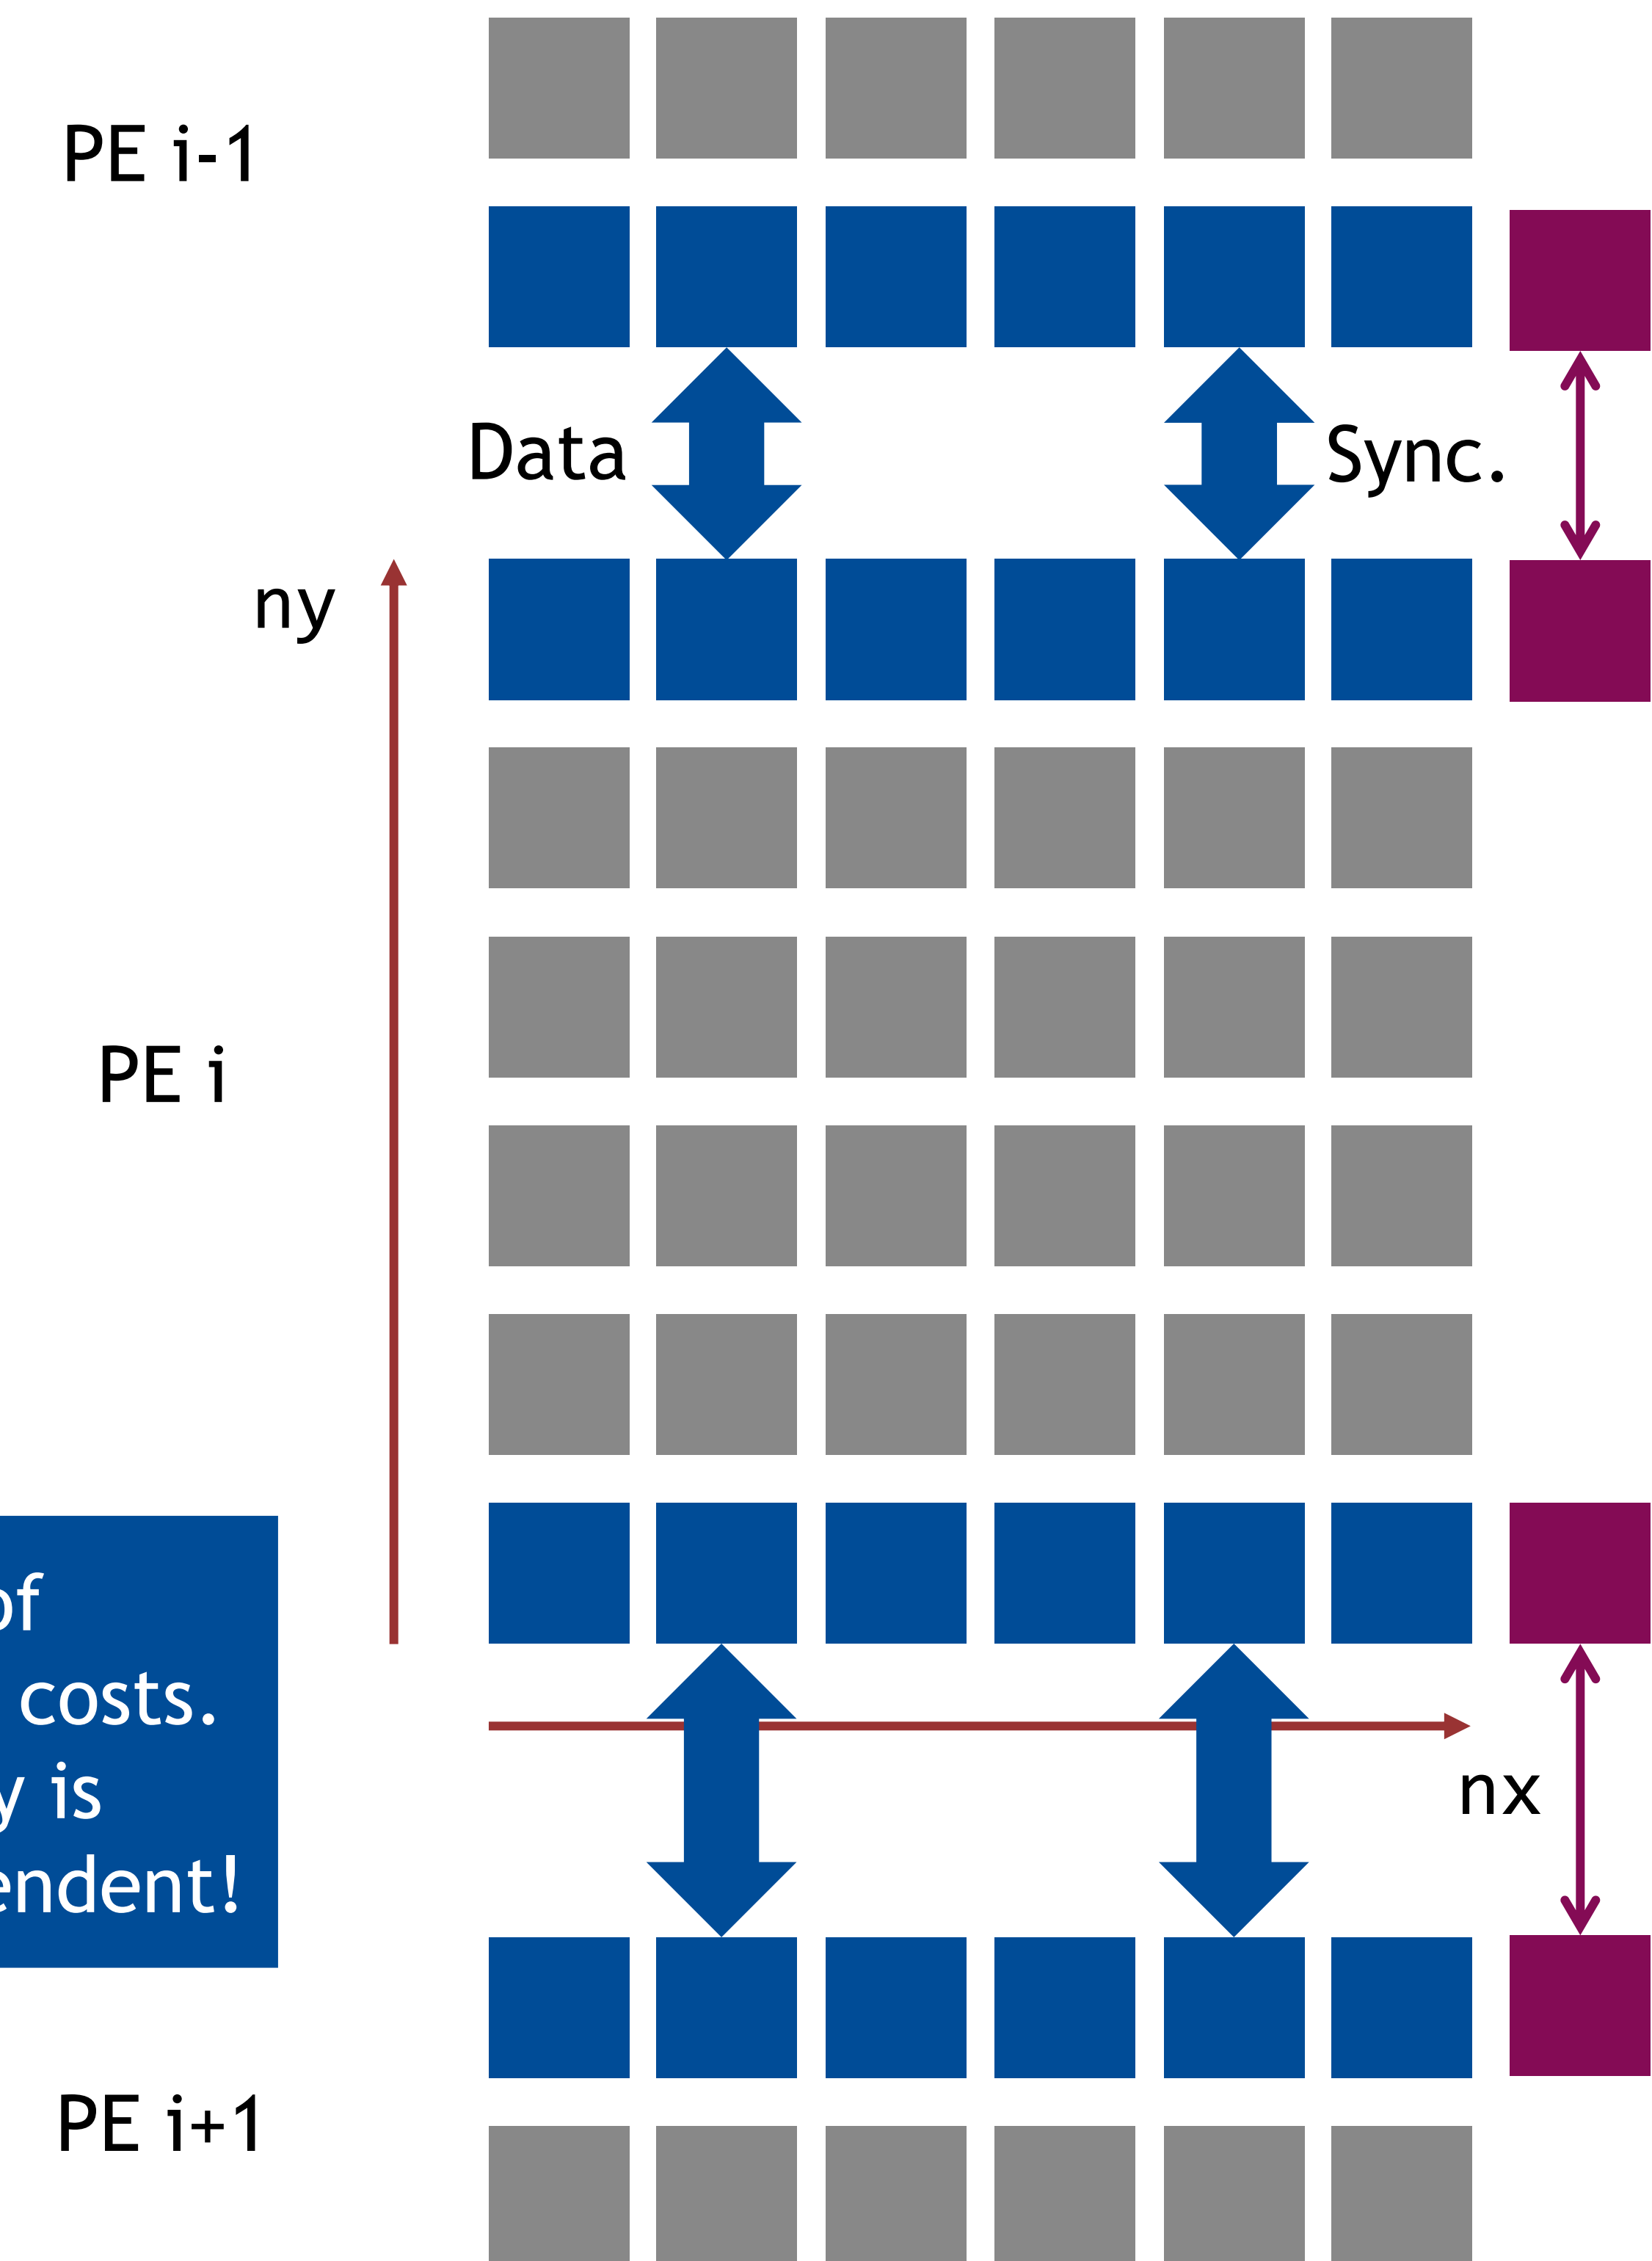
<https://docs.nvidia.com/hpc-sdk/nvshmem/api/docs/gen/api/rma.html#stdrmatypes>

In-Kernel Synchronization

- Point-to-point synchronization across PEs within a kernel
- Enables kernel fusion

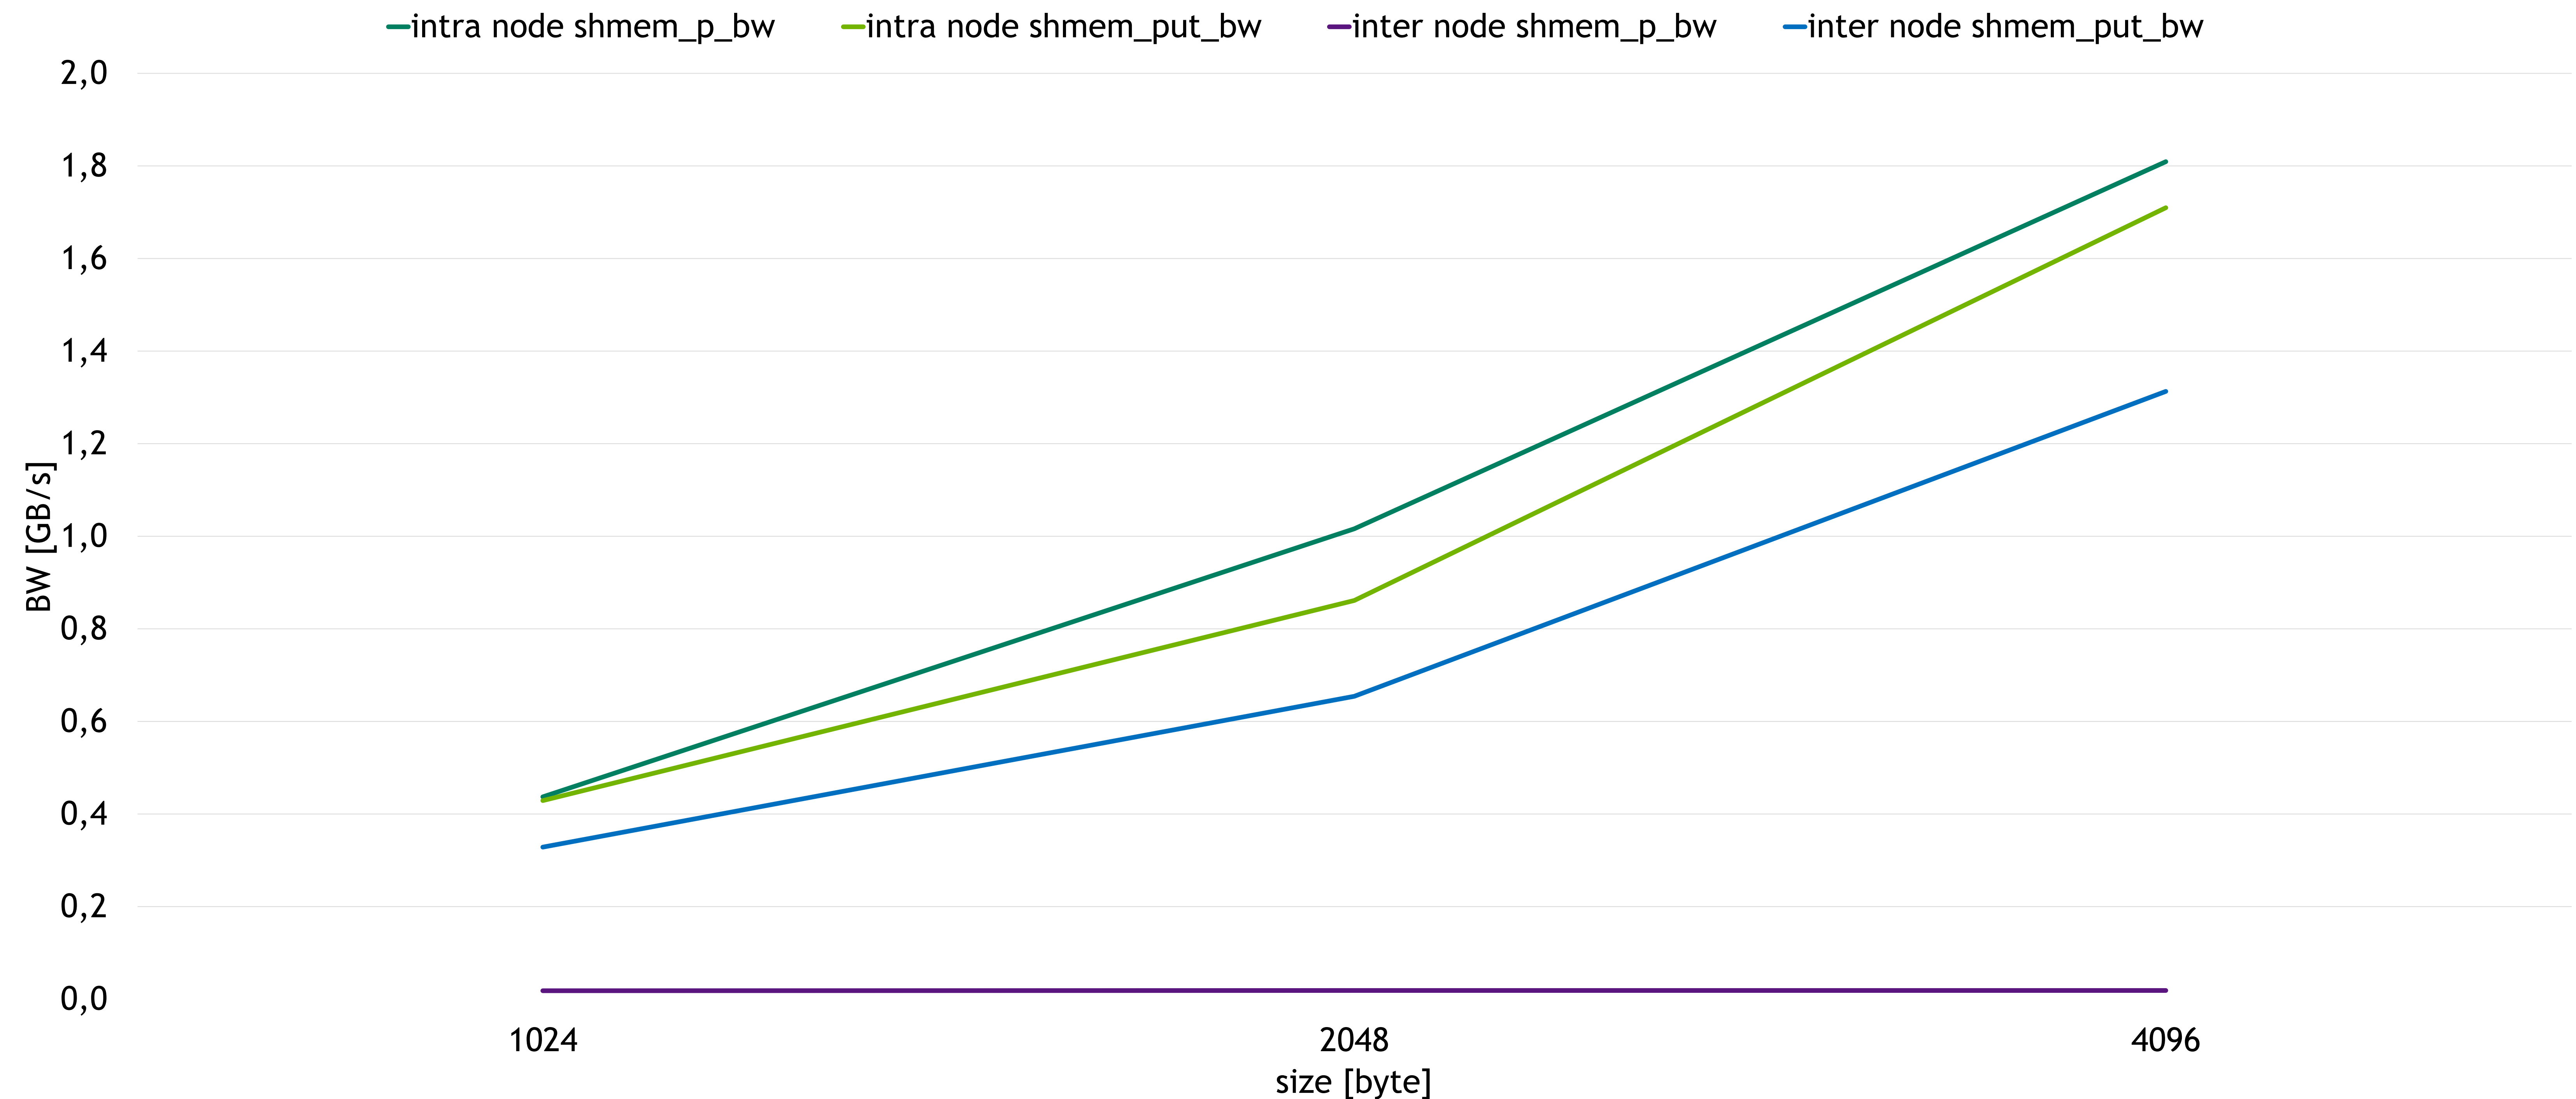
```
__global__ void stencil_multi_step(float *u, float *v, int N, int *sync, ...) {
    int ix = get_ix(blockIdx blockIdx.threadIdx);
    int iy = get_iy(blockIdx blockIdx.threadIdx);
    for (int iter = 0; iter < N; iter++) {
        swap(u, v); compute(u, v, ix, iy);
        // Thread block-level data exchange (assume even/odd iter buffering)
        int boffset = get_block_offset(blockIdx blockIdx);
        if (blockIdx.y == 0)
            nvshmemx_float_put_nbi_block(u+(ny+1)*nx+boffset, u+nx+boffset, blockDim.x, top_pe);
        if (blockIdx.y == (blockDim.y-1))
            nvshmemx_float_put_nbi_block(u + boffset, u+ny*nx+boffset, blockDim.x, bottom_pe);
        if (blockIdx.y == 0 || blockIdx.y == (blockDim.y-1)) {
            __syncthreads();
            nvshmem_quiet();
            if (threadIdx.x == 0 && threadIdx.y == 0) {
                nvshmem_atomic_inc(sync, top_pe);
                nvshmem_atomic_inc(sync, bottom_pe);
            }
        }
        nvshmem_wait_until(sync, NVSHMEM_CMP_GT, 2*iter*gridDim.x);
    }
}
```

Be aware of
synchronization costs.
Best strategy is
application dependent!




NVSHMEM perftestS

shmem_p_bw and shmem_put_bw on JUWELS Booster - NVIDIA A100 40 GB



NVSHMEM perftestS

shmem_p_bw and shmem_put_bw on JUWELS Booster - NVIDIA A100 40 GB




Optimized Intra-Node Communication

Supported on NVLink and PCI-E

Use CUDA IPC or cuMem* API to map symmetric memory of intra-node PEs into virtual address space

`nvshmem_[put|get]` on device -> load/store

`nvshmem_[put|get]_on_stream` -> `cudaMemcpyAsync`



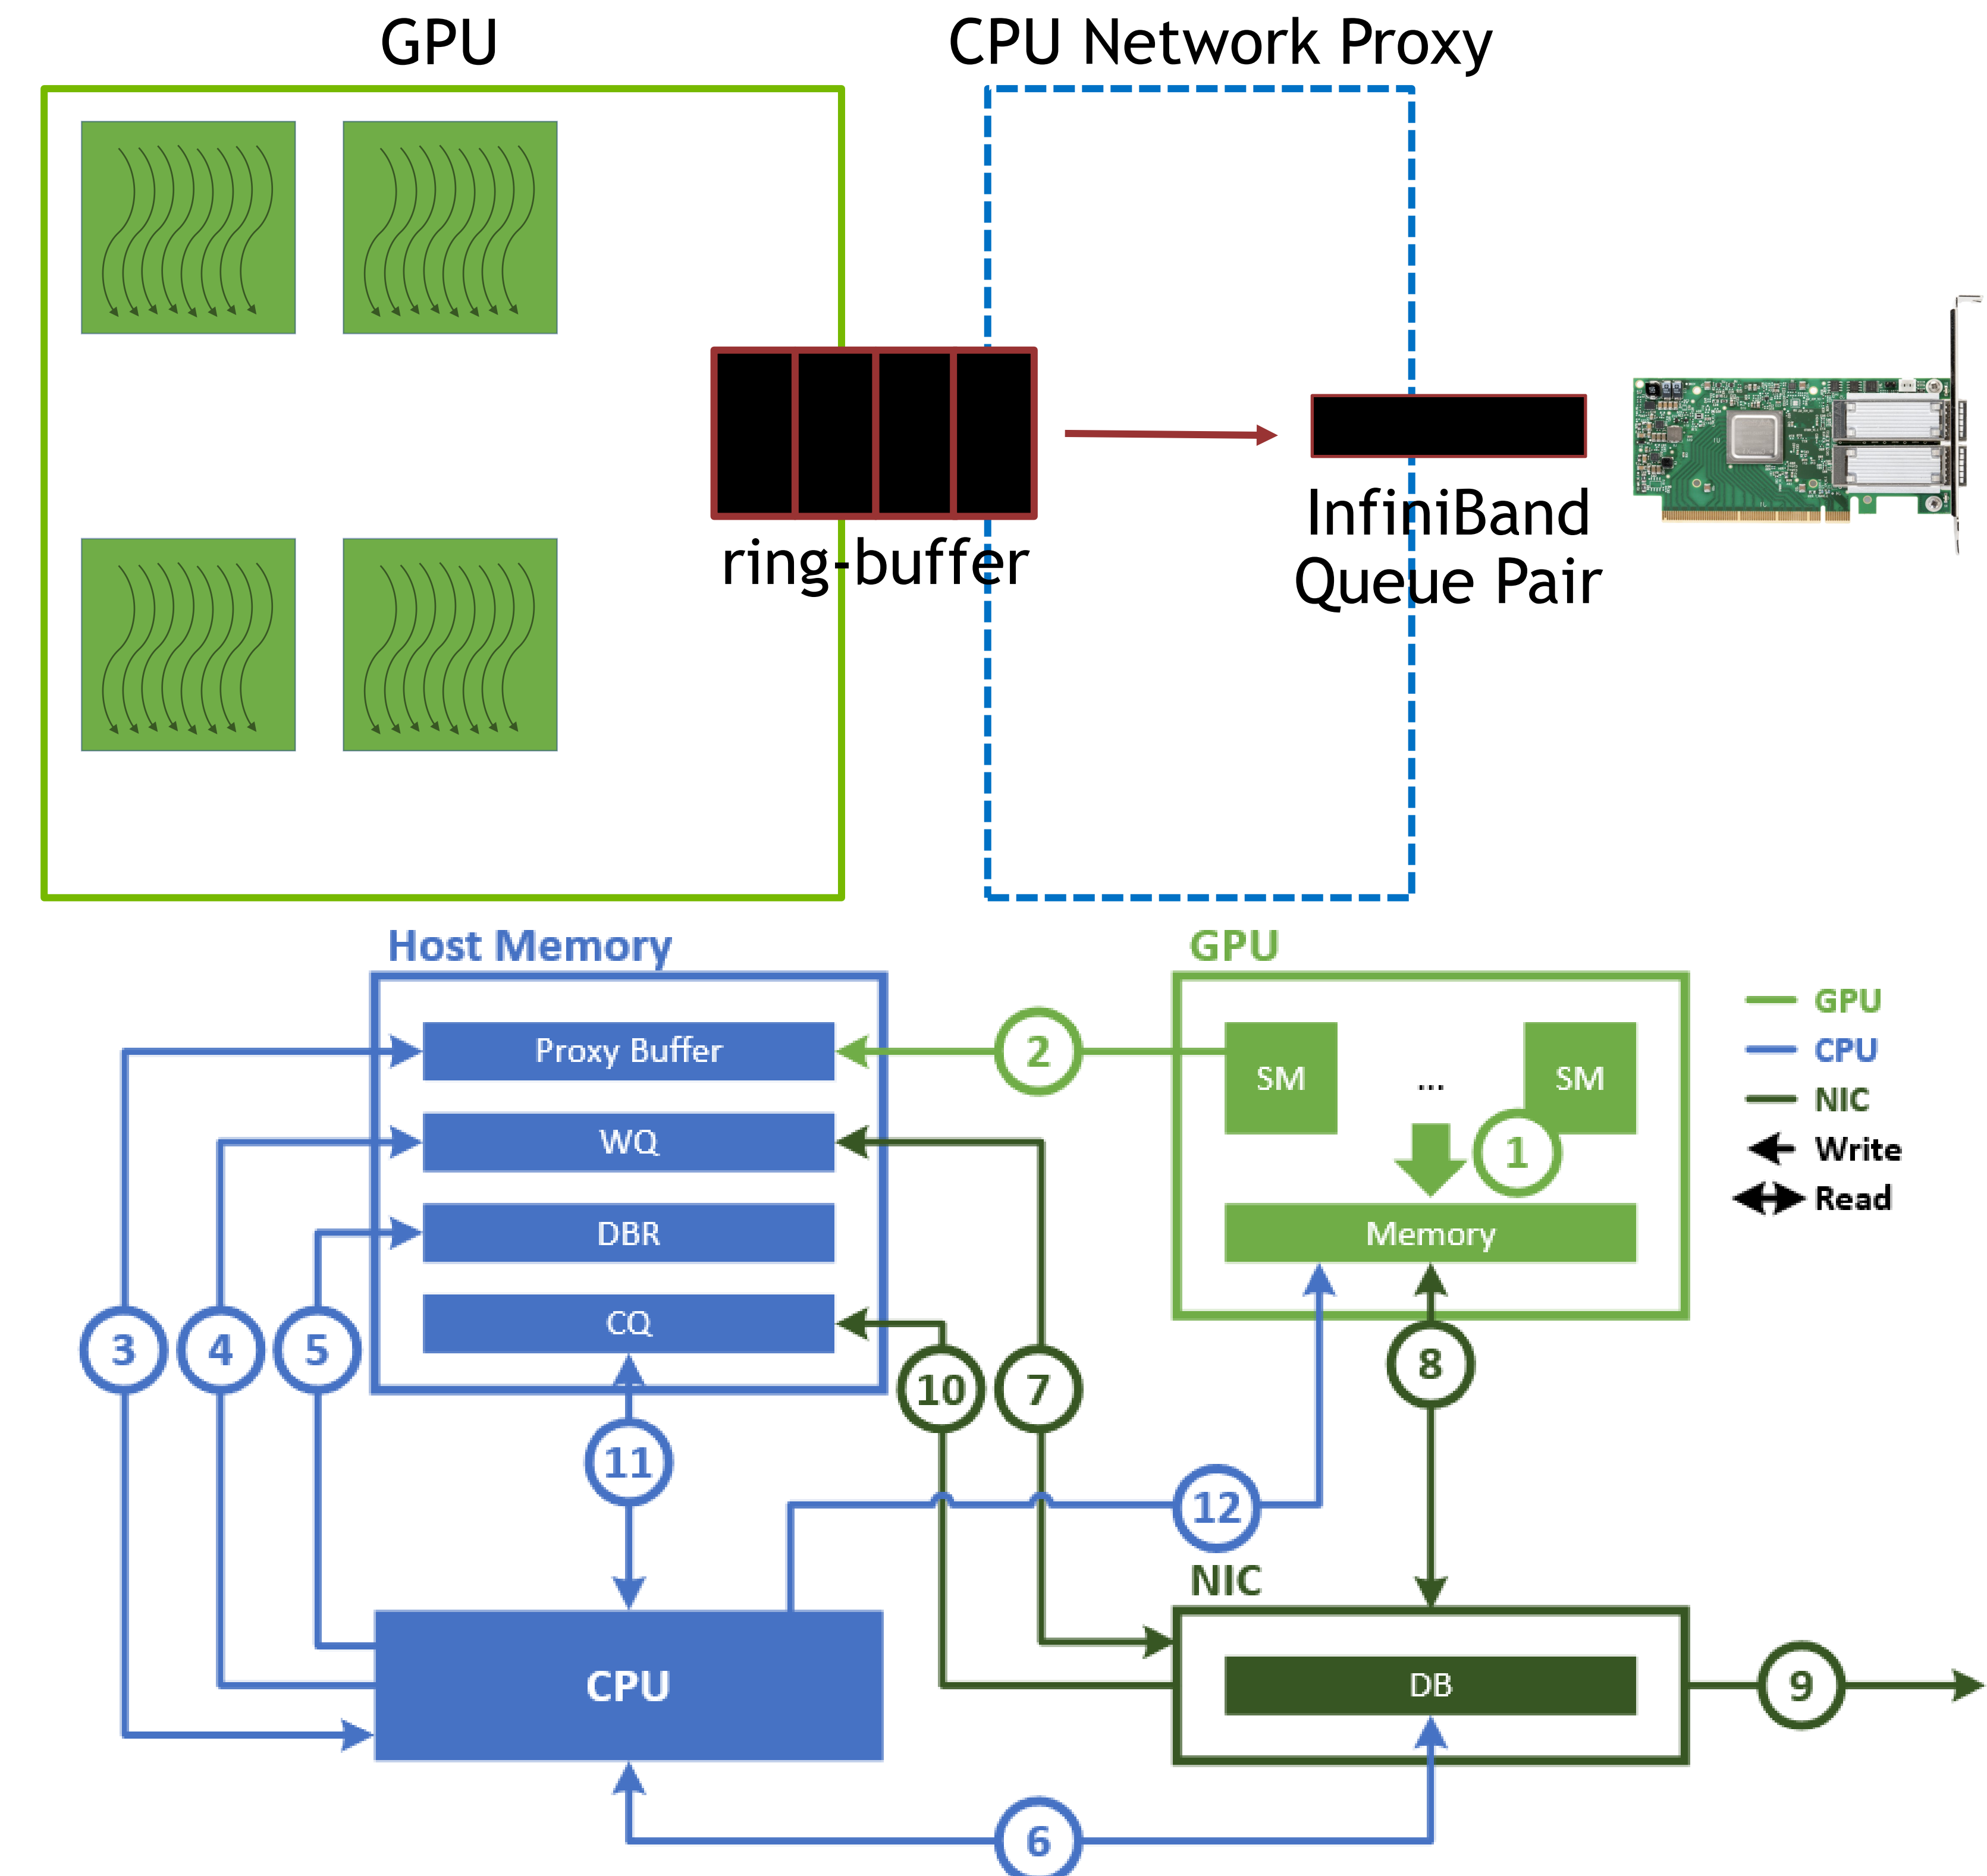
Optimized Inter-Node Communication

NVSHMEM supports inter-node communication over InfiniBand, RoCE, and UCX (experimental)

Using GPUDirect RDMA (data plane)

Reverse offloads network transfers from GPU to the CPU (control plane)

Ring buffer implementation avoids memory fences when interacting with CPU network proxy

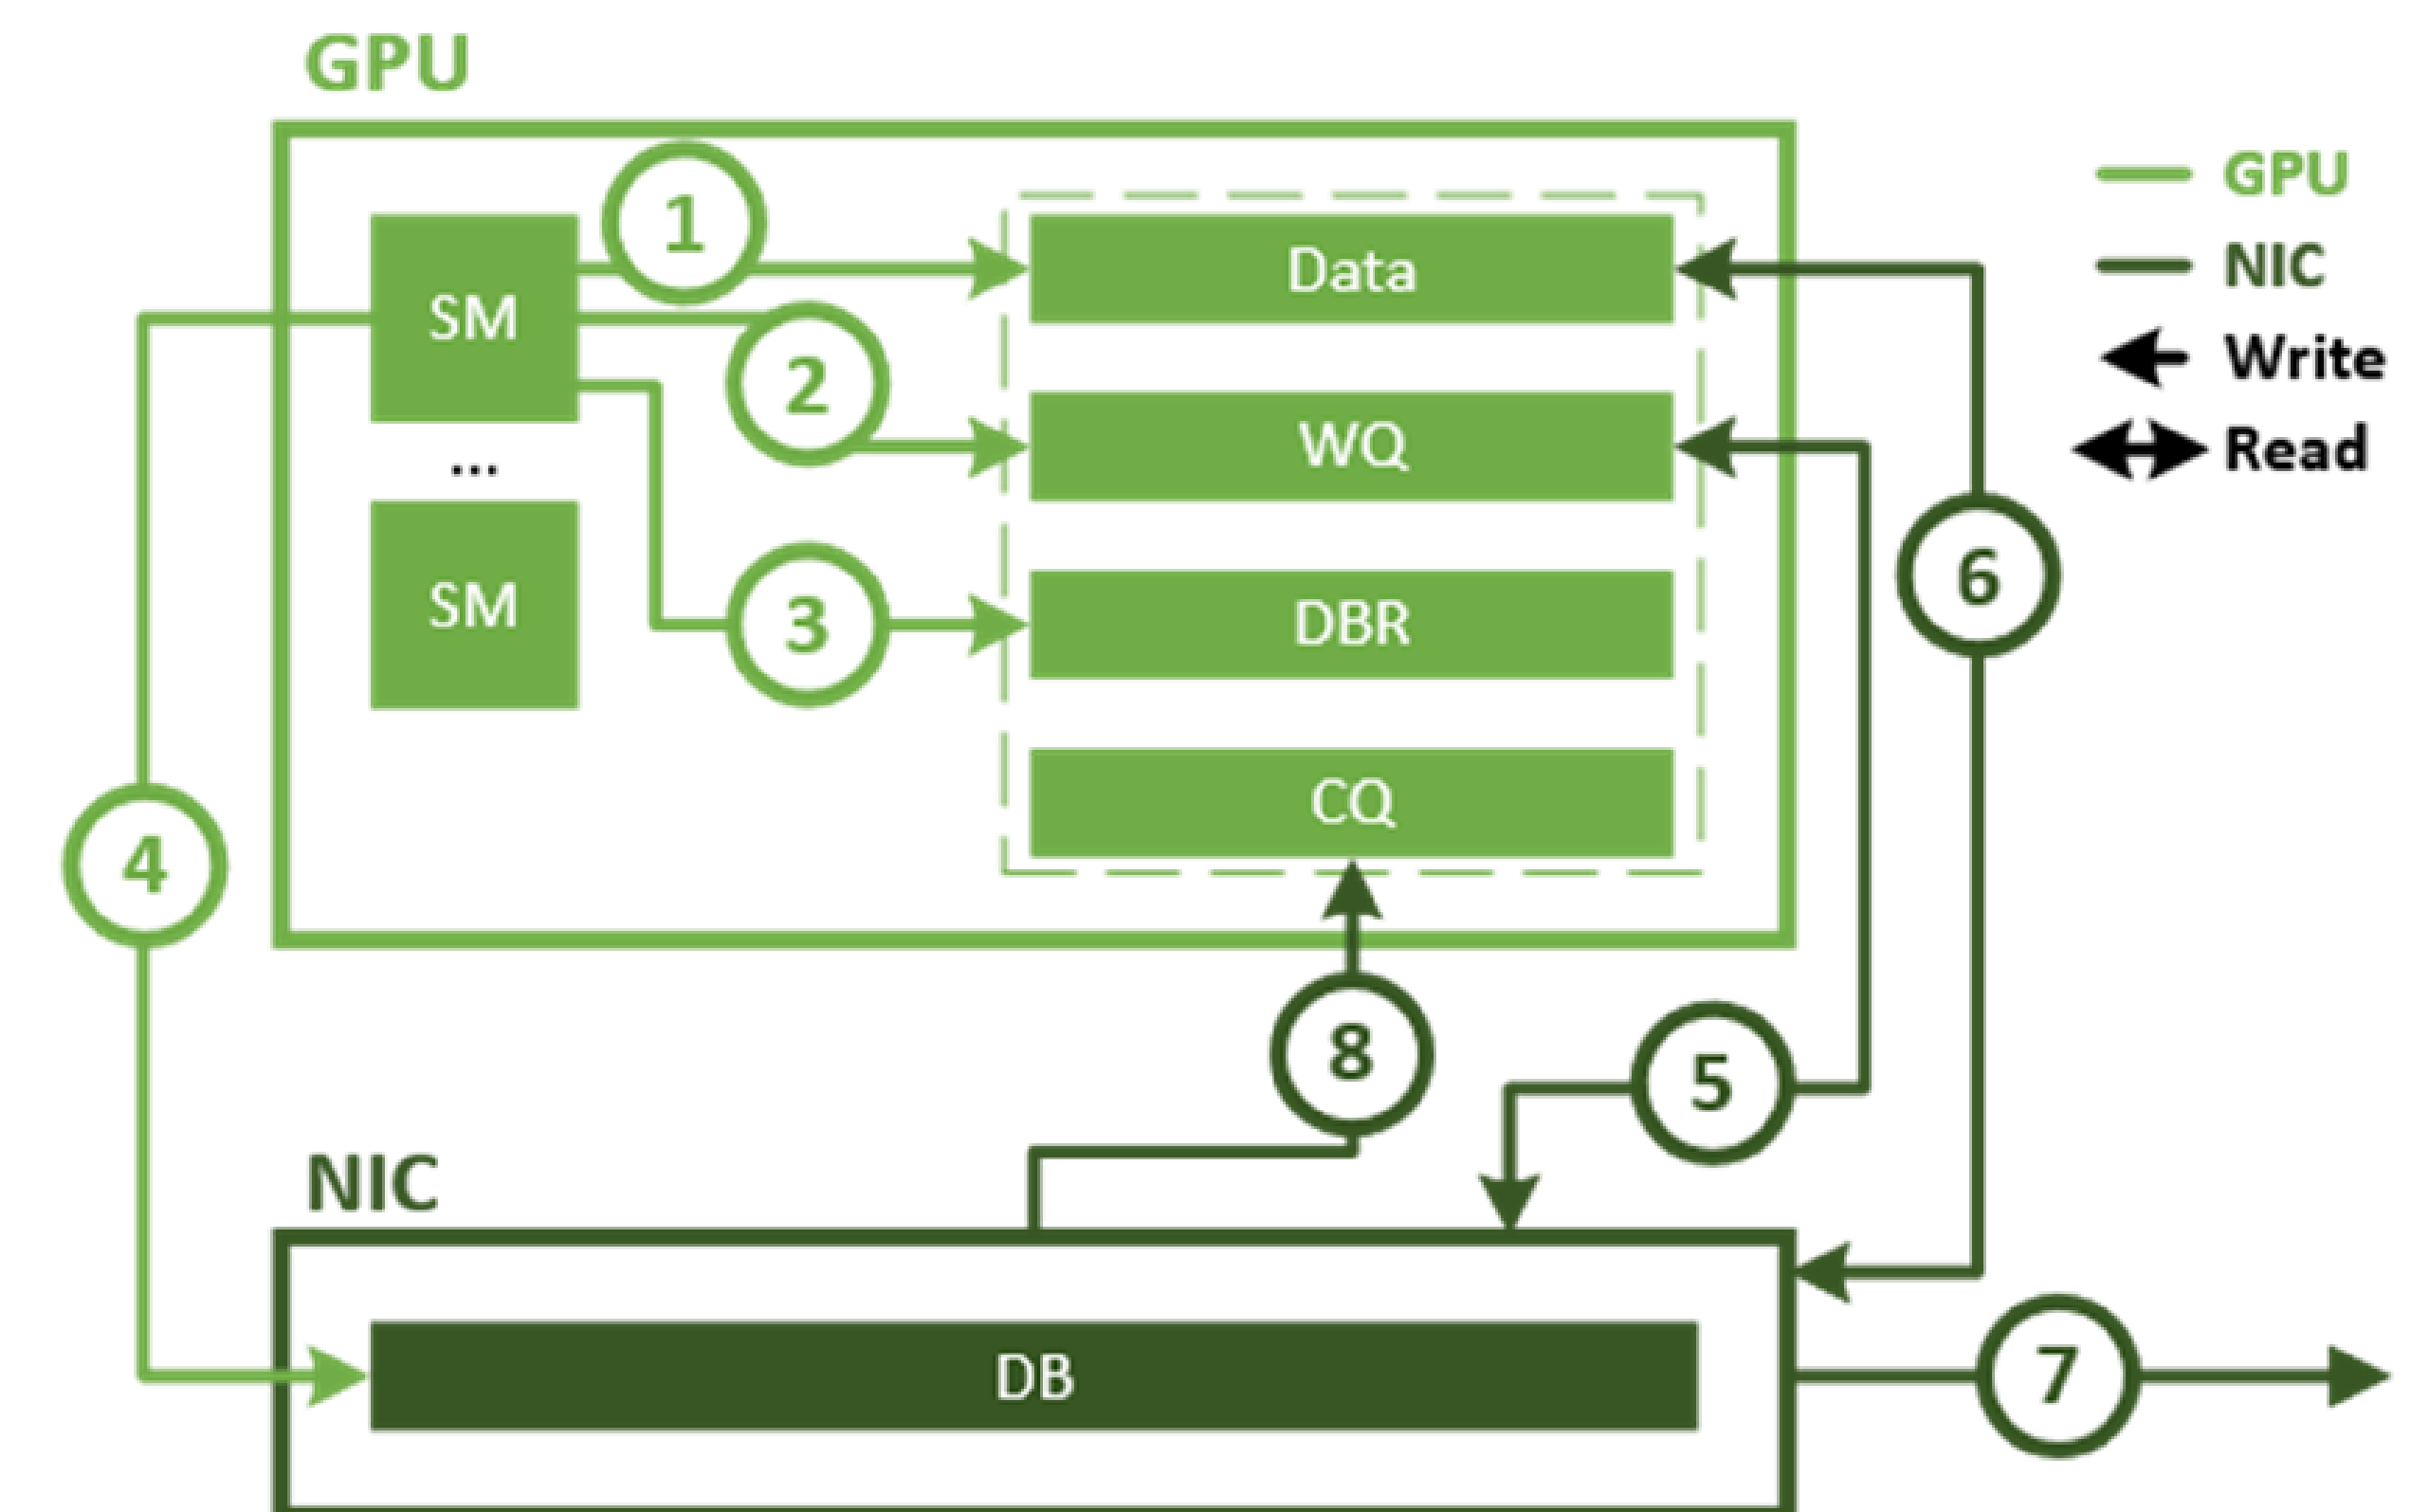
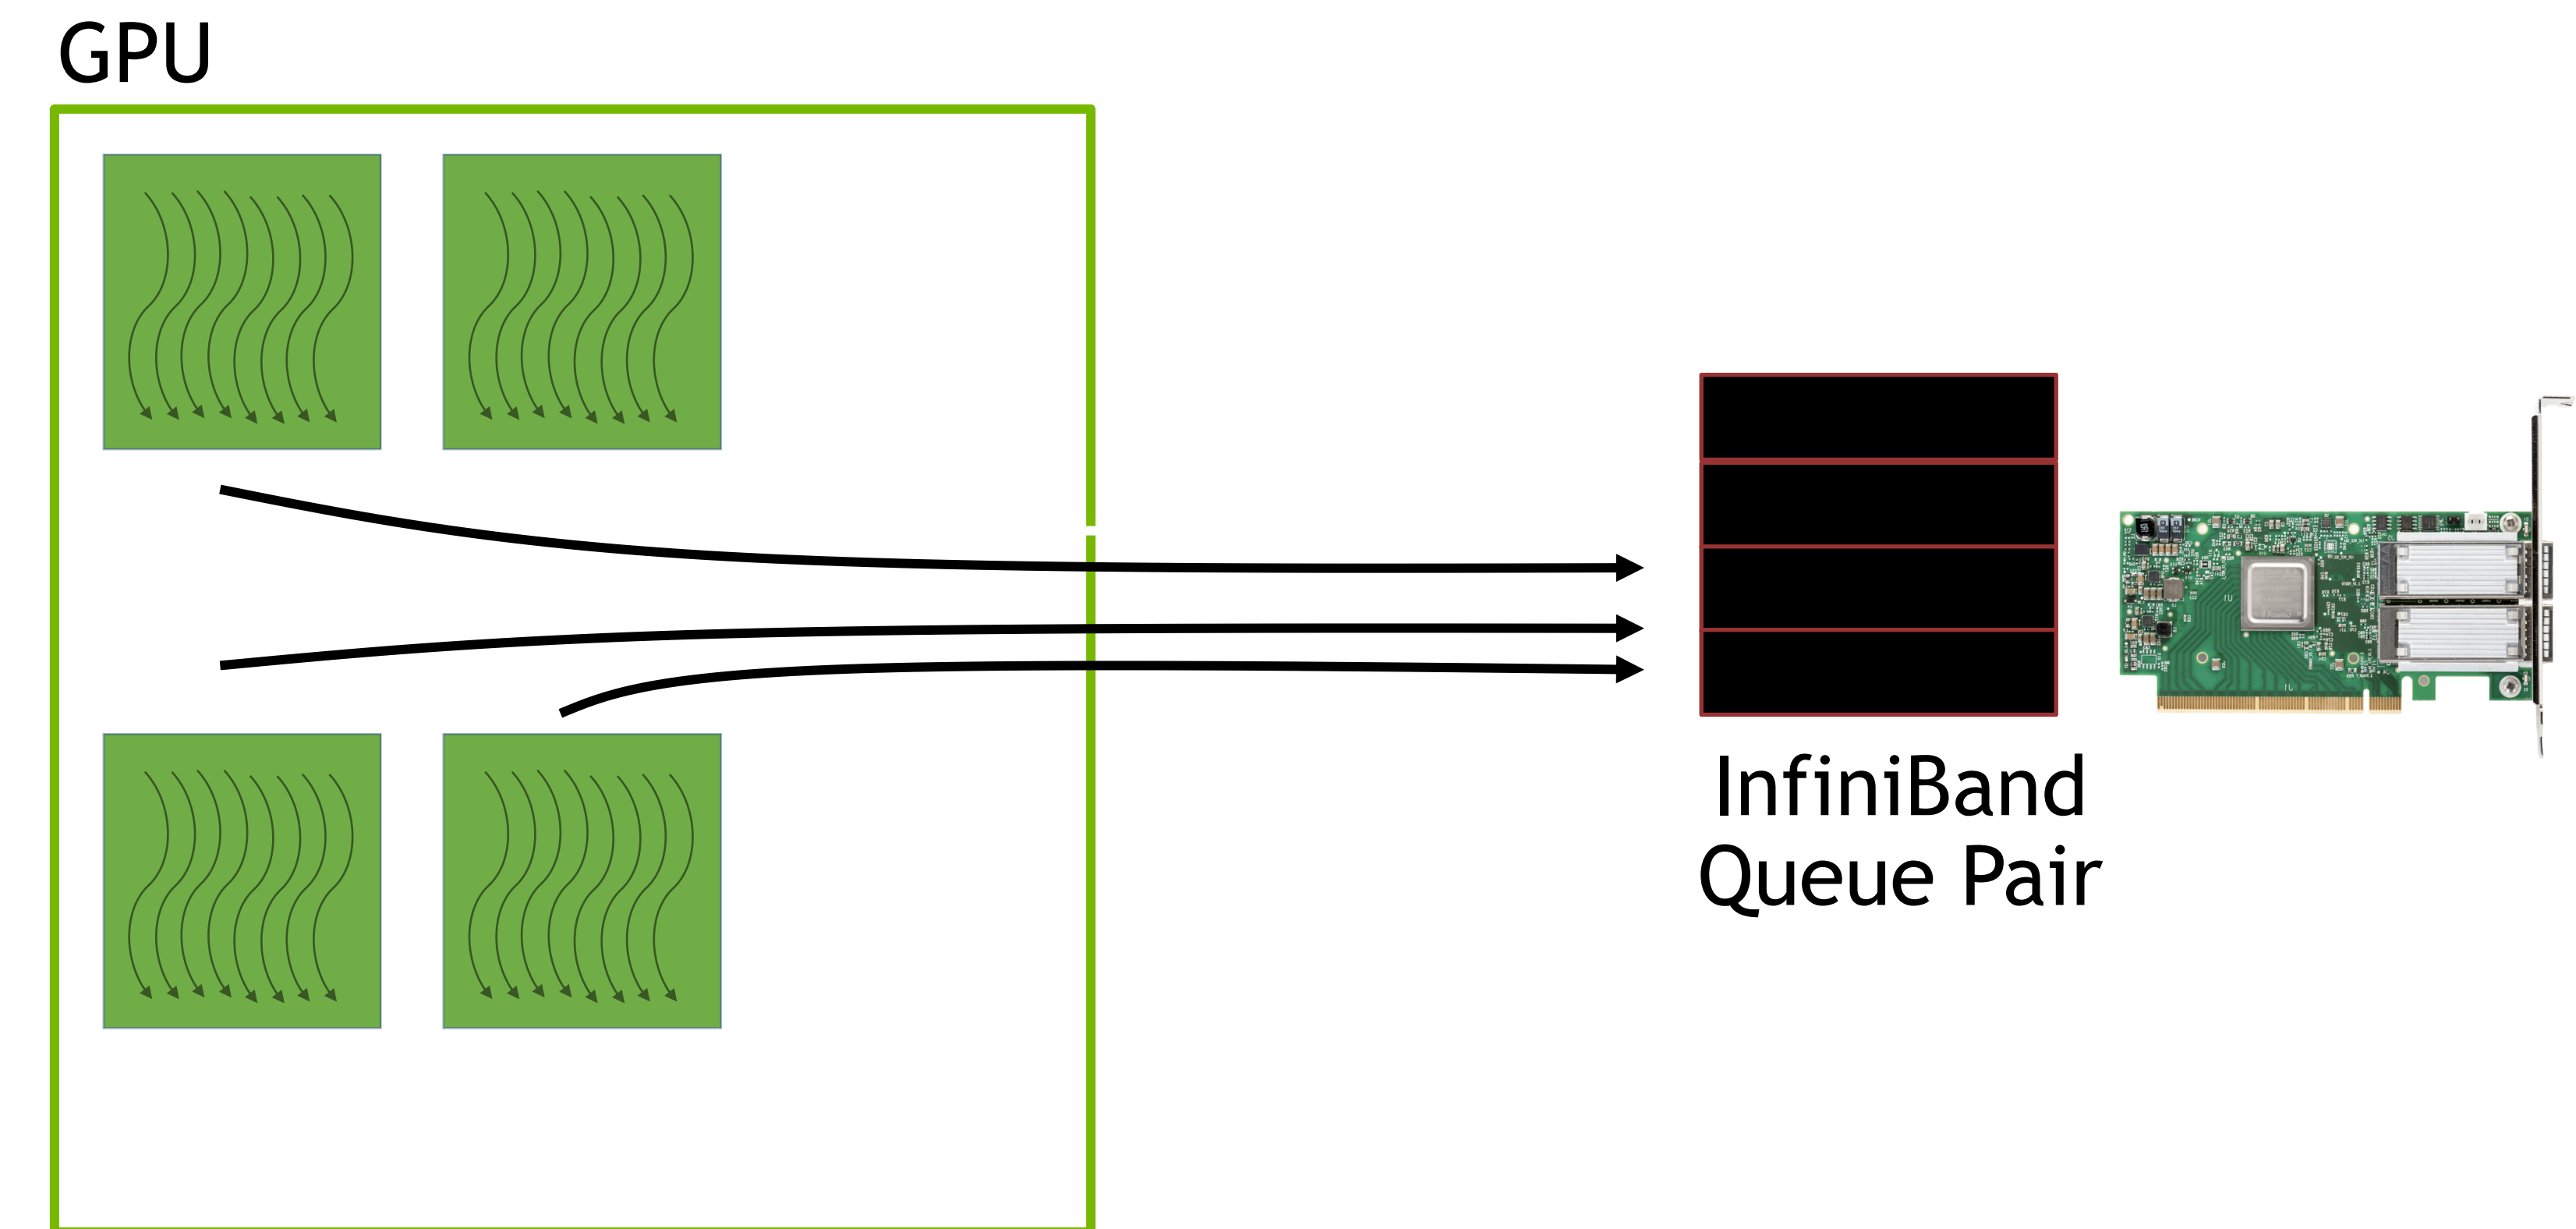


Optimized Inter-Node Communication Improved

IB GPUDirect Async (IBGDA) over InfiniBand


Using GPUDirect RDMA (data plane)

GPU directly initiates **network transfers** involving the CPU only for the setup of control data structures




NVSHMEM perftestS with IBGDA

shmem_p_bw and shmem_put_bw on JUWELS Booster - NVIDIA A100 40 GB




NVSHMEM Version with NVL, CPU Proxy and IBGDA



QUDA Strong Scaling on Selene

Lattice Quantum ChromoDynamics



<https://www.nvidia.com/en-us/on-demand/session/gtcsj20-s21673/>

up to 1.6x Speedup (512 GPUs half precision)

- CUDA Graphs help minimize CPU-side launch overhead and Device-side execution overhead
- Device-initiate communication enables:
 - fine grained communication and computation overlap with sometimes less coding effort
 - kernel fusion not possible with host initiate communication models like MPI and NCCL
- With **IB GPUDirect Async** (IBGDA) NVSHMEM can achieve peak Network message rates
- Without IBGDA for good intranode device-initiated communication performance it is necessary to aggregate larger messages (nvshmemx_TYPENAME_put_nbi_block)

- CUDA Graphs documentation <https://docs.nvidia.com/cuda/cuda-c-programming-guide/index.html#cuda-graphs>
- NVSHMEM: CUDA-Integrated Communication for NVIDIA GPUs (a Magnum IO session): <https://www.nvidia.com/en-us/on-demand/session/gtcspring22-s41044/>

Overcoming Latency Barriers: Strong Scaling HPC Applications with NVSHMEM:

- <https://www.nvidia.com/en-us/on-demand/session/gtcsj20-s21673/>
- <https://developer.nvidia.com/blog/scaling-scientific-computing-with-nvshmem/>
- <https://developer.nvidia.com/blog/accelerating-nvshmem-2-0-team-based-collectives-using-nccl/>